

Database Scalability



What we will learn today ?

1. Recap

2. Database
Replication Models

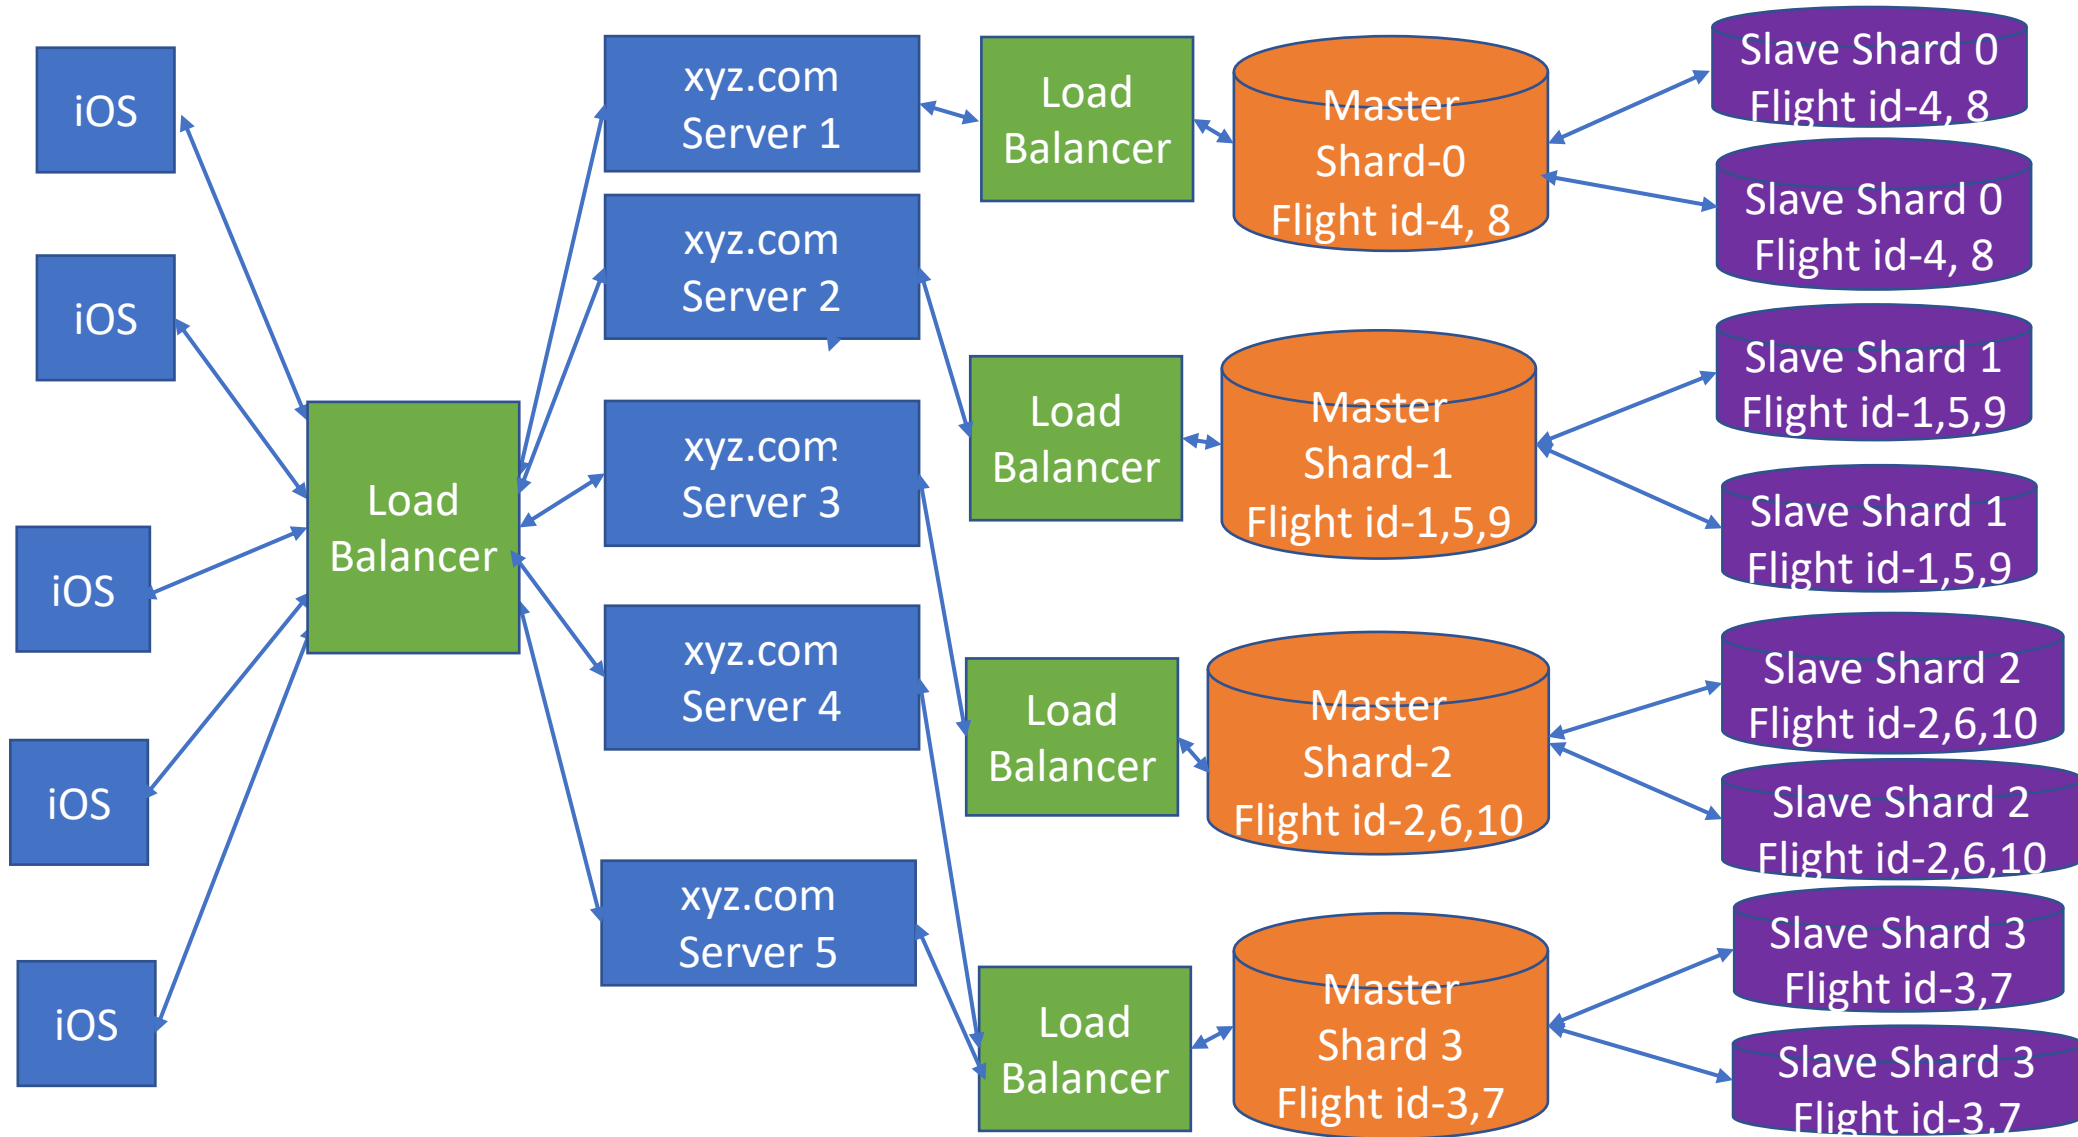
3. 2 Phase Commit

4. CAP



1. Recap

1. Recap: Database Replication



1. Recap: Database Replication

1. Now that we have **3 dB replicas for each Shard** what happens if we update one shard Replica ?
2. How do we make sure that **updates are consistent across** shard replicas ?
3. How do we make sure **read/write operations are in sync** across all the 3 shard Replicas ?
4. How do the **dB replicas communicate** among themselves ?
5. We have few **dB replication models(Master-Slave, Master-Master)** to to fix this problem.



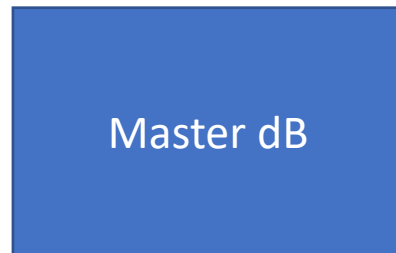
2. Database Replication Models

2. Database Replication Models

1. Single Primary Replication Model(Master-Slave)
2. Multi Primary Replication Model(Master-Master)

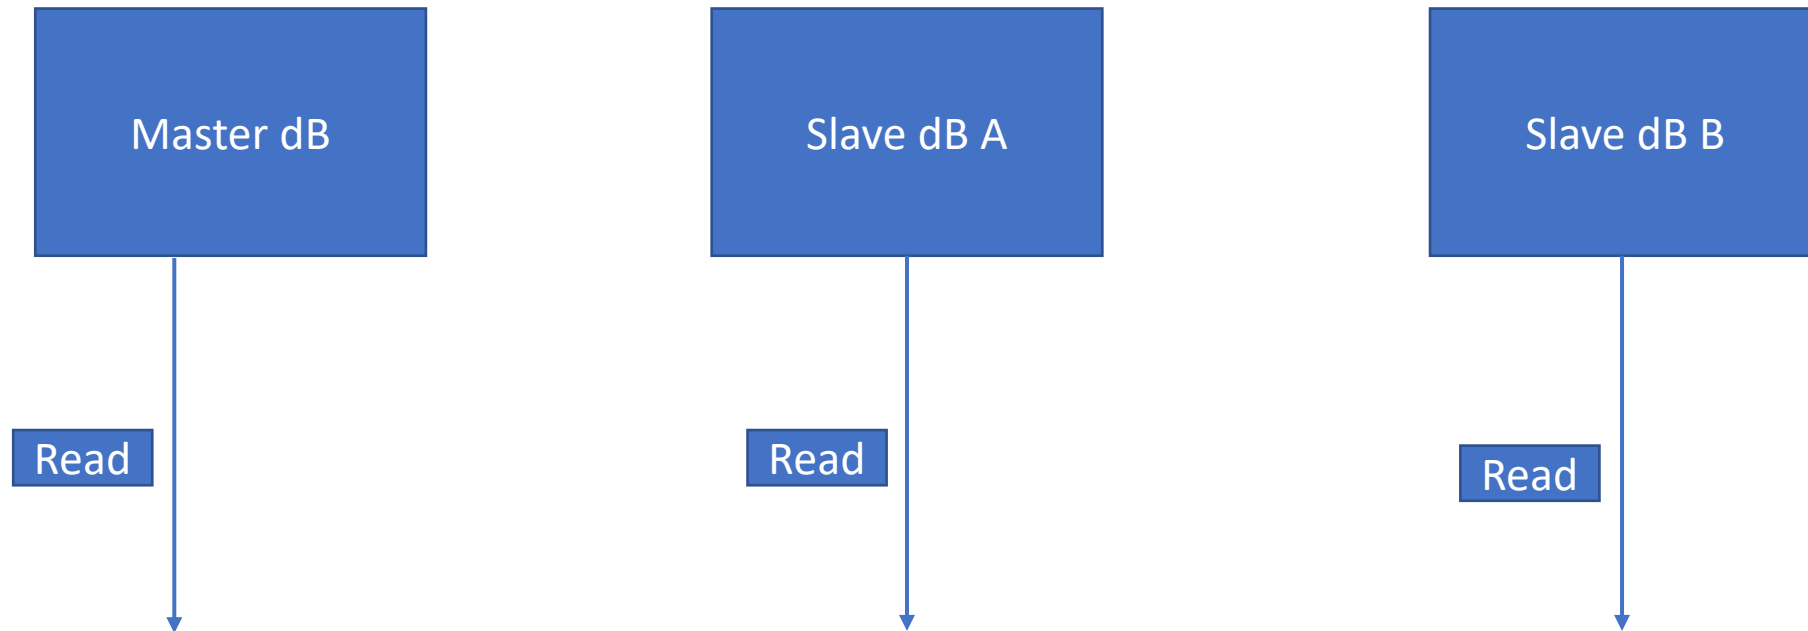
2.1 Single Primary Replication Model

1. This model helps in **distributing the read** request loads across dBs.
2. In this model we have **two types of dBs- A Primary dB(Master) and Secondary dBs(Slave)**.
3. **Primary dB(Master)**: We can perform **read/write** operation on this dB.
4. **Secondary dBs(Slaves)**: We can perform **ONLY read** operation on this dB.



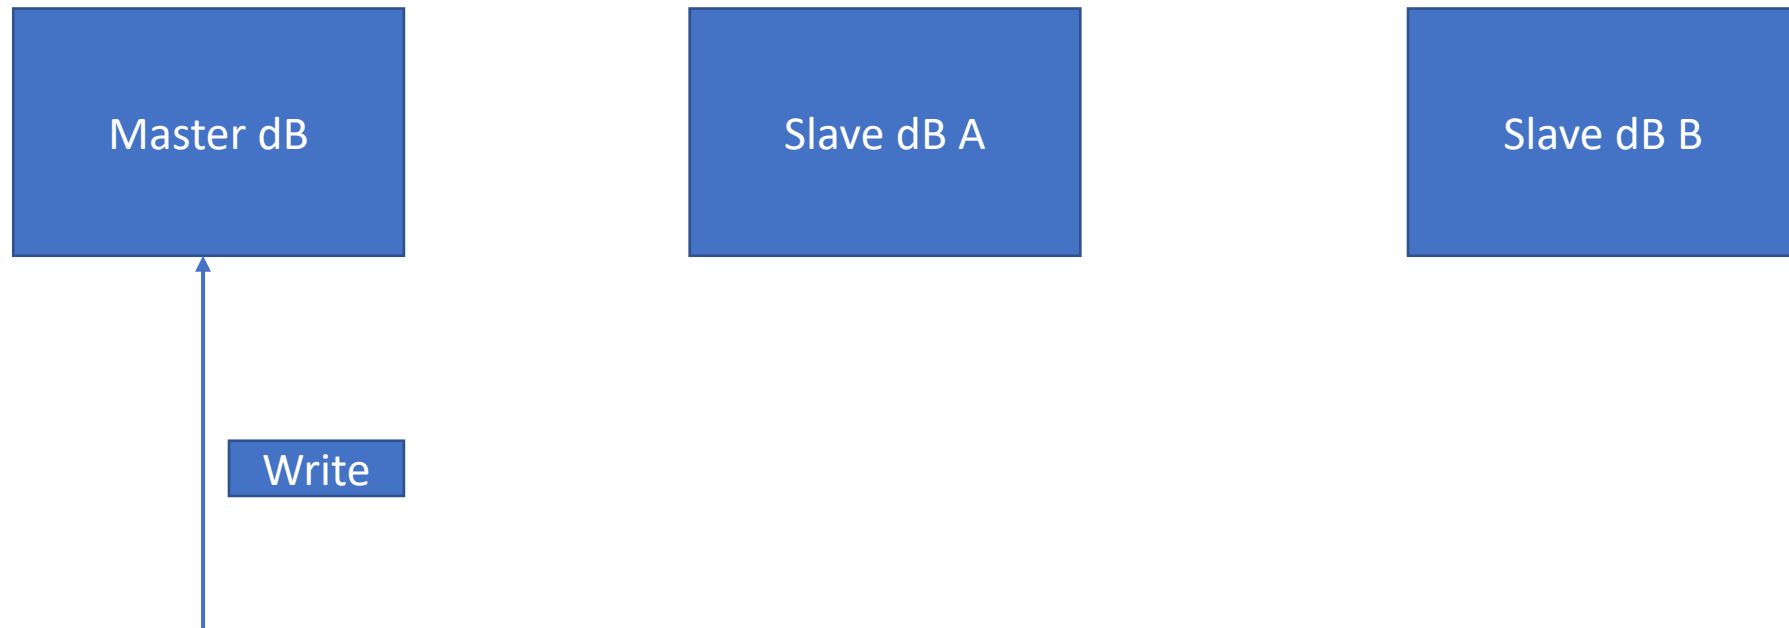
2.1 Single Primary Replication Model

1. Every **Read request** can be redirected to any of the 3 dB Replica.
2. We use a load balancer to distribute read request load evenly.



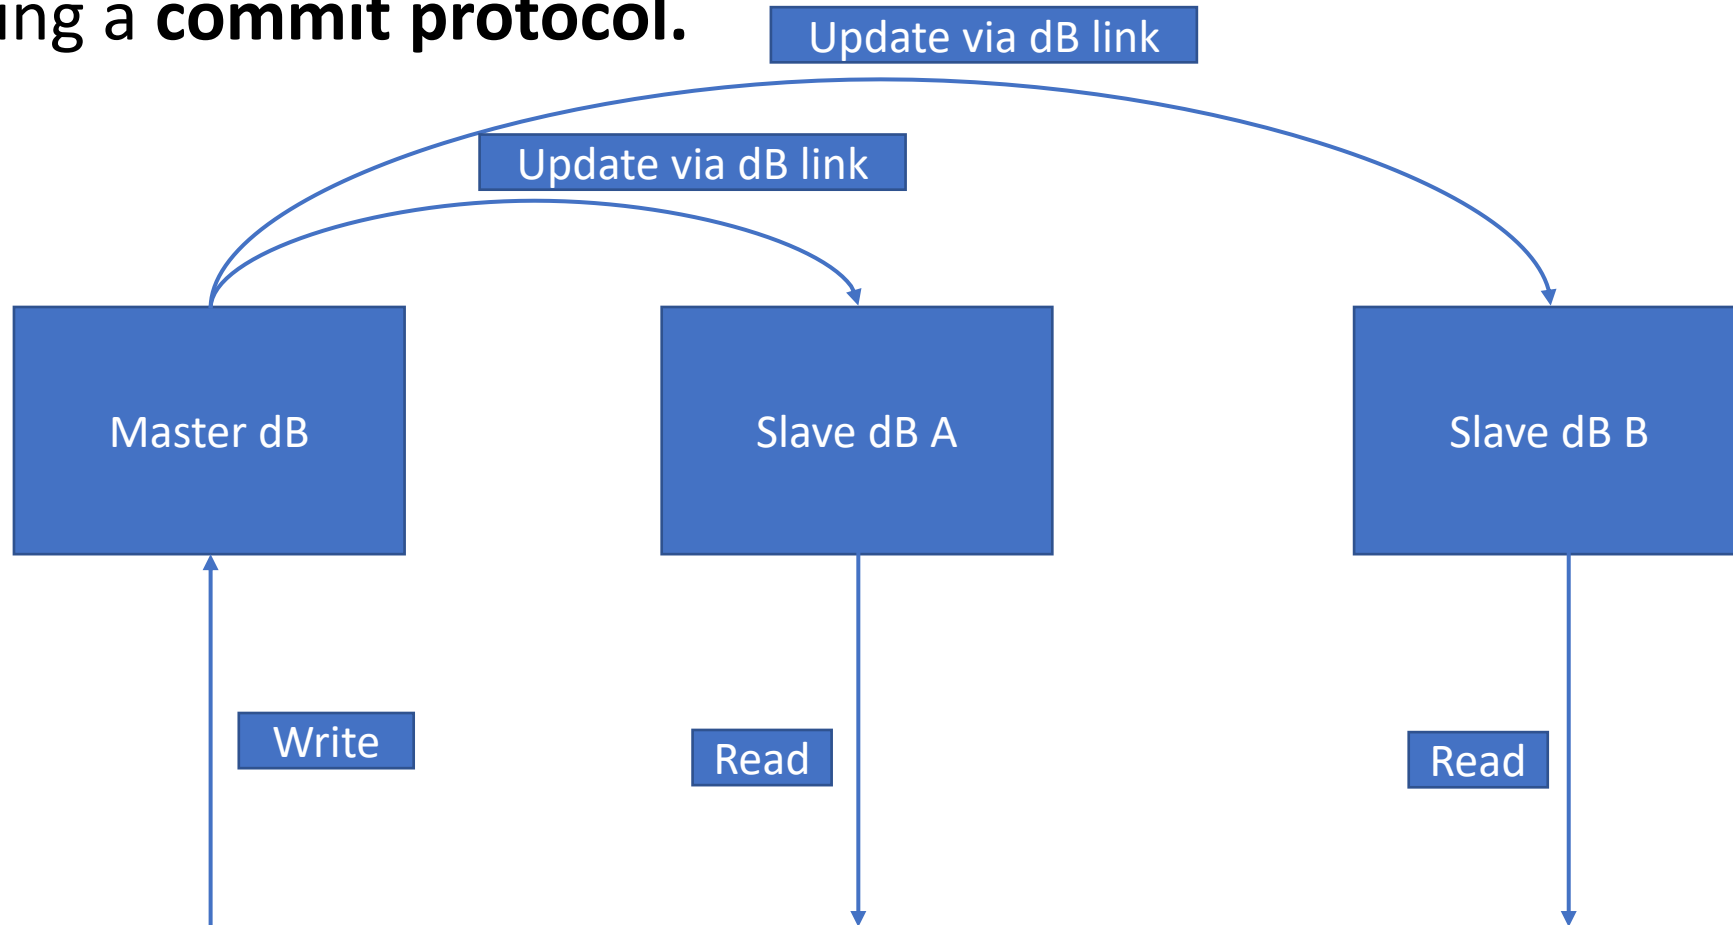
2.1 Single Primary Replication Model

What happens if a write to master(Primary) dB is made ?



2.1 Single Primary Replication Model

Master(Primary) dB performs a sync/async update to its slaves via dB links using a **commit protocol**.

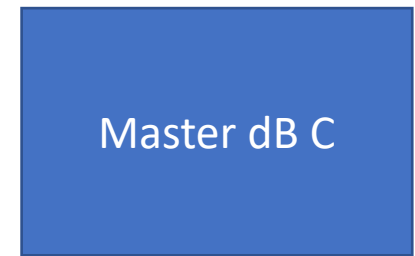
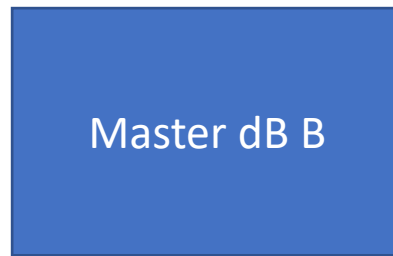
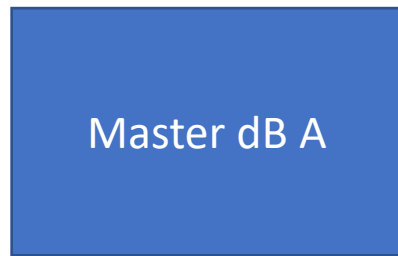


2.1 Concerns with Single Primary Replication Model

1. We might run into **race condition**: If we **write** some data into **primary(master)** dB and then try to **read** data from a **secondary(slave)** dB before the **primary(master)** dB updates its changes to the **secondary(slave)** dB.
2. This is a good model for a **news website** where **Reads > Writes**.
3. If the **primary(master)** dB fails, we cant write to the dB. SPOF.
4. How do we **distribute the write requests** across dBs ? (write heavy)

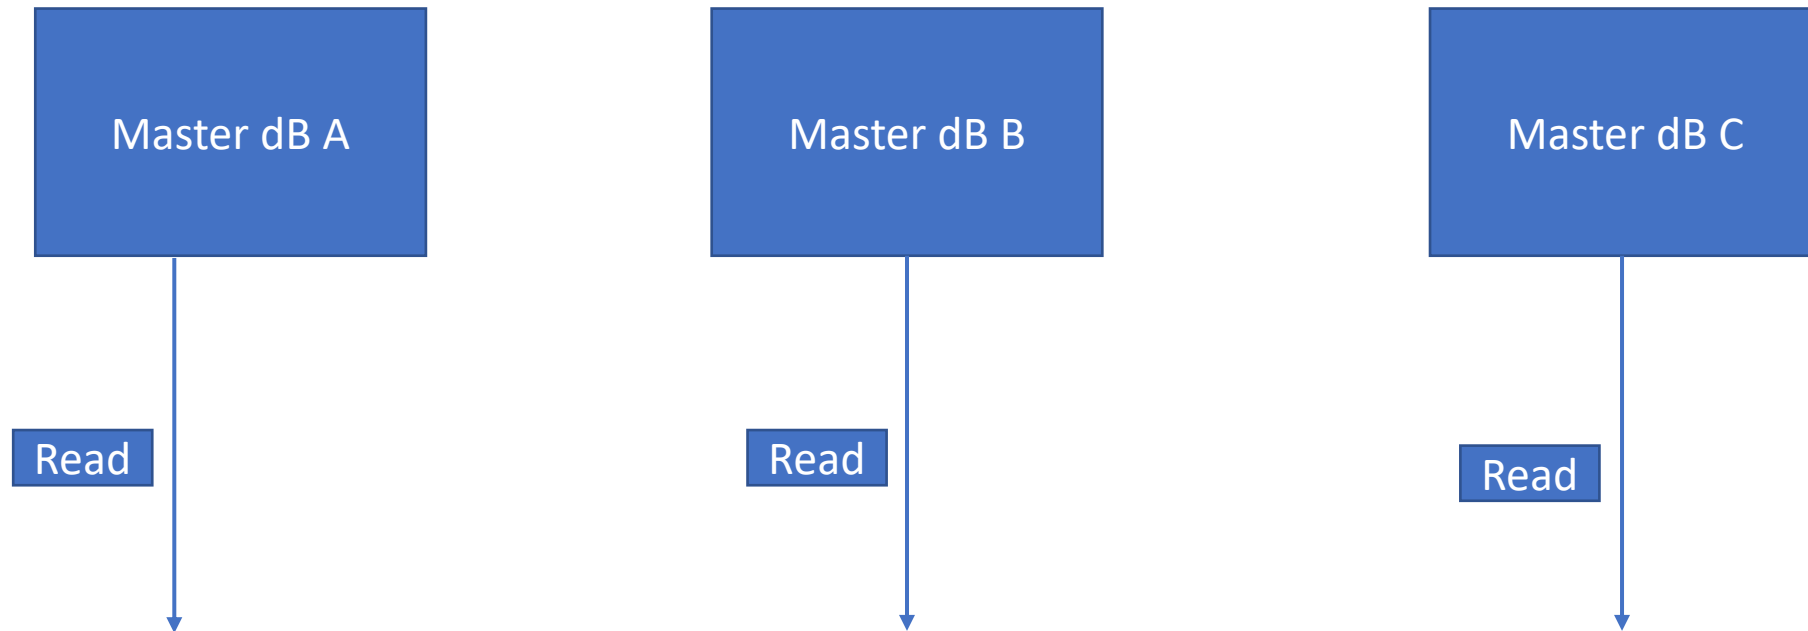
2.2 Multi Primary Replication Model

1. This model helps in **distributing the read+write** request loads across dB servers.
2. We have **multiple primary(master)** dB to read from and write to the dBs.



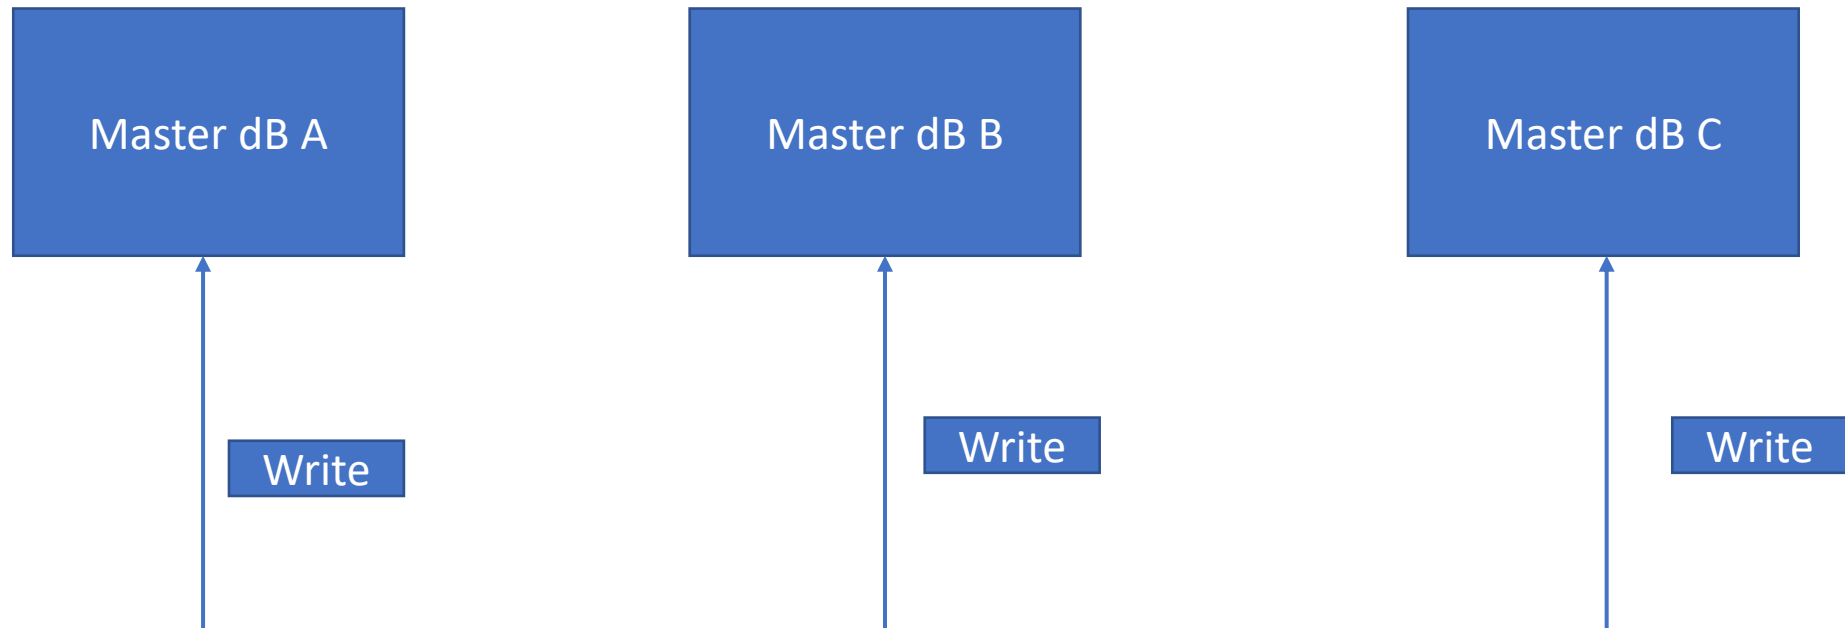
2.2 Multi Primary Replication Model

1. Every **Read request** can be redirected to any of the 3 dB Replica.
2. We use a load balancer to distribute read request load evenly.



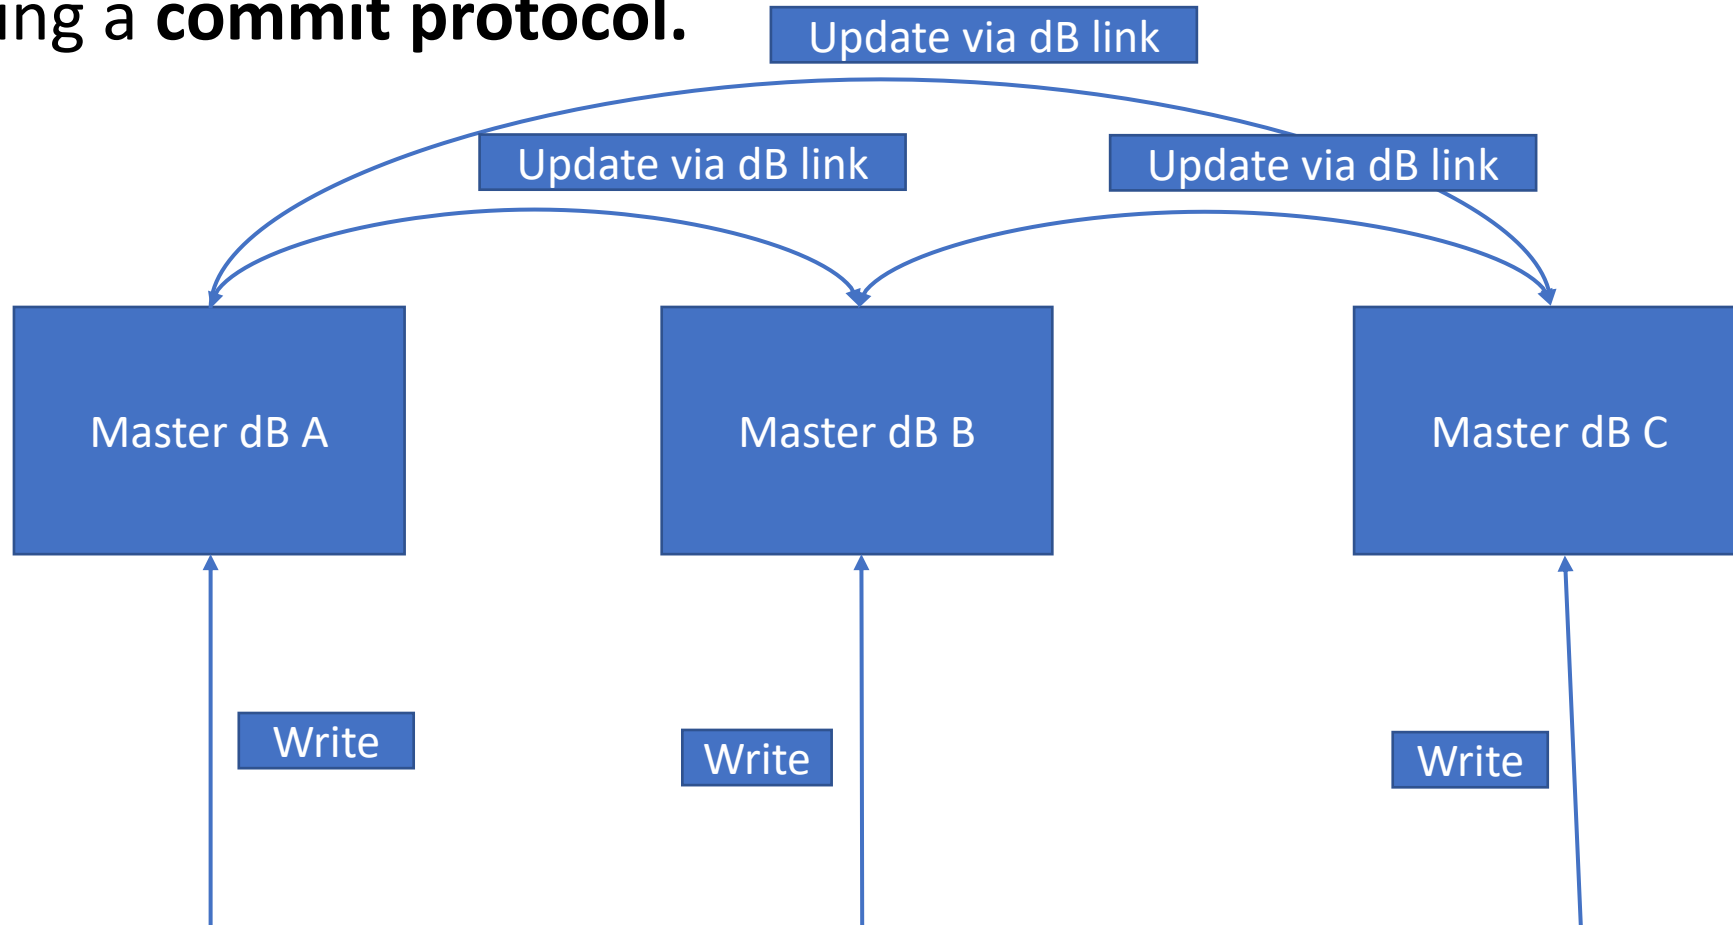
2.2 Multi Primary Replication Model

1. Every **Write request** can be redirected to any of the 3 dB Replica.
2. We use a load balancer to **distribute write request** load evenly.



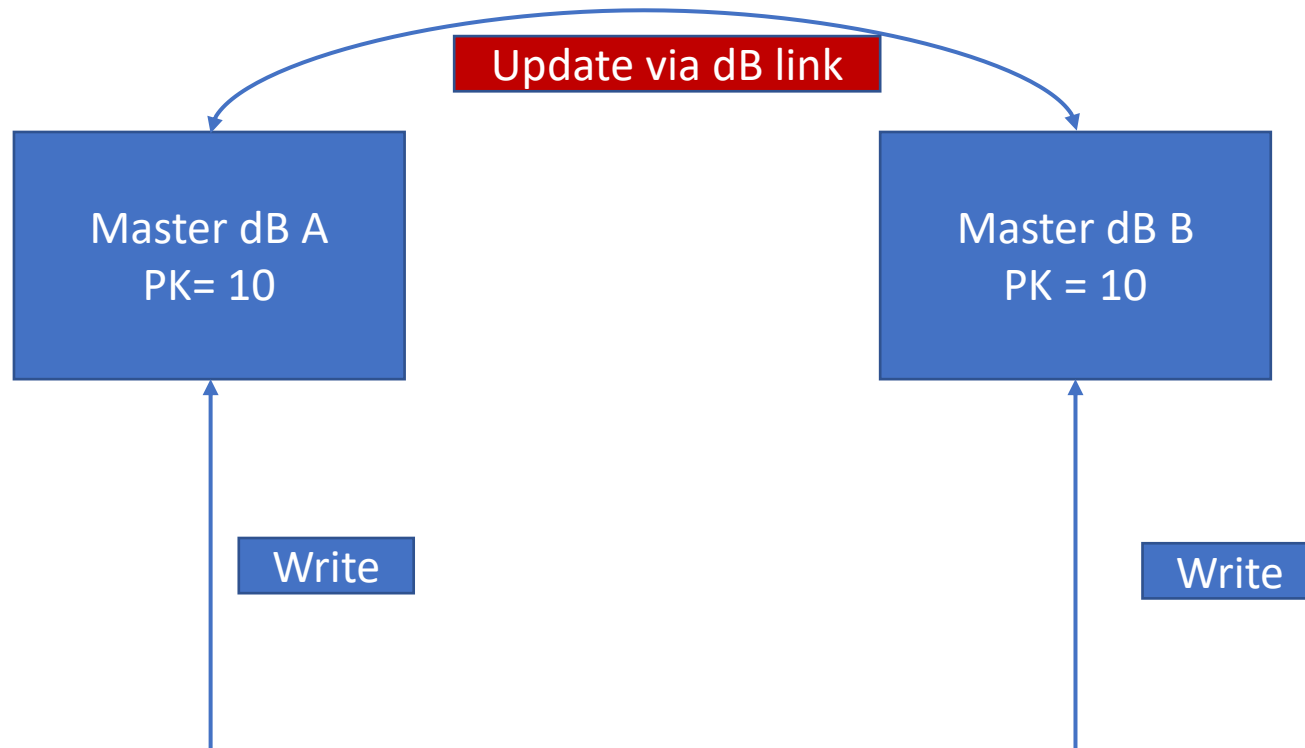
2.2 Multi Primary Replication Model

Master(Primary) dB performs a sync/async update other Master via dB links using a **commit protocol**.



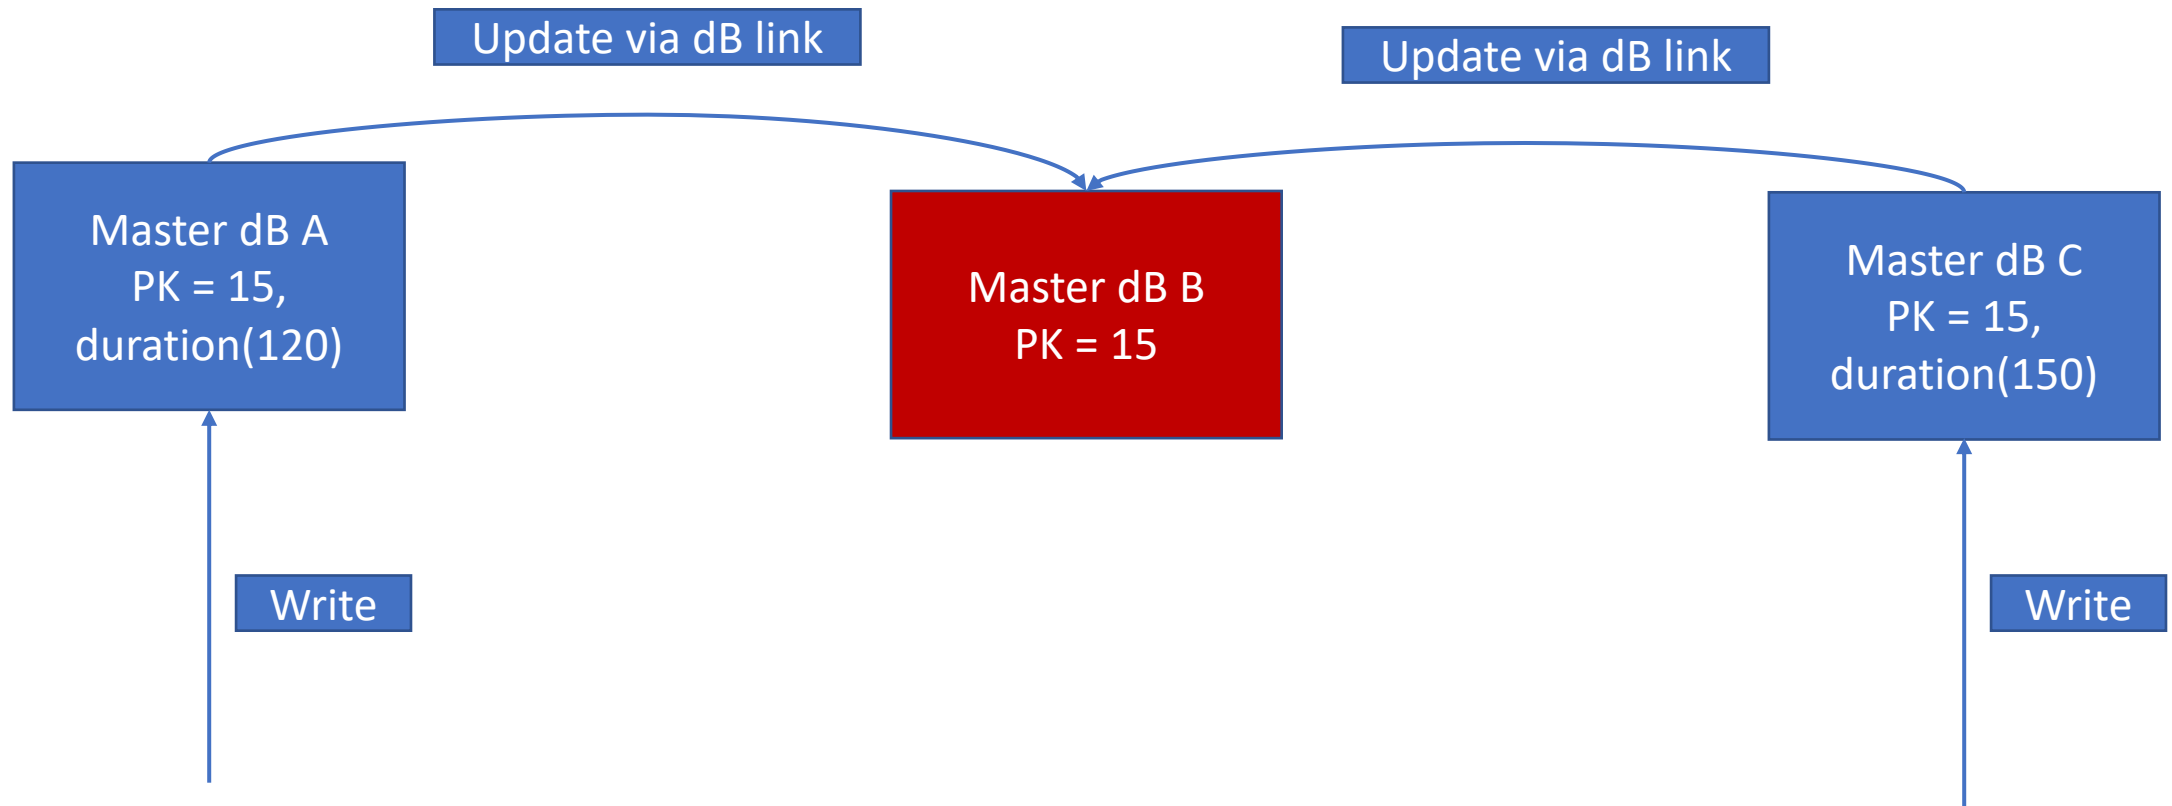
2.2 Concerns with Multi Primary Replication Model

1. We might run into **Primary Key Conflict**: If we **create** a new row into each of two **primary(master)** dBs A and B having the **same primary key**, and then an update happens between A and B to sync data.



2.2 Concerns with Multi Primary Replication Model

1. What happens when two Master dBs A and C try to update the same row in Master dB B ? Conflict





3. 2-Phase Commit Protocol

3. 2-Phase Commit Protocol

1. Generally a **dB transaction on a single dB server** follows either of two semantics to maintain its data Consistency:

ACID = **A**tomicity + **C**onsistency + **I**solation + **D**urability(SQL dBs)

BASE = **B**asically Available + **S**oft State + **E**ventual Consistency(NoSql dbs like Key-Value store, document store, wide-Column store, graph dB)

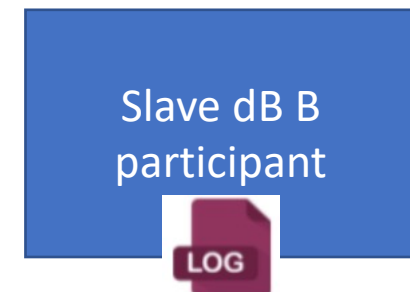
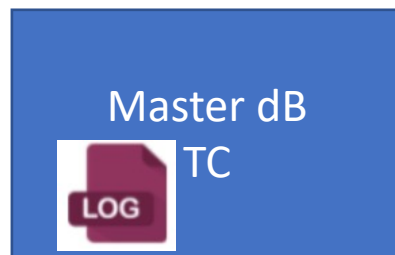
2. How do we **maintain ACID/BASE property** in case of **distributed dB servers** ?

3. 2-Phase Commit Protocol

1. This is a **distributed atomic(all-or-nothing) commit protocol** to achieve **agreement(consensus) across dB replicas** during a commit process.
2. It has **two phases**:
 - A. Commit Request Phase(Voting phase)
 - B. Commit Phase
3. Lets see how it works with a running example of a **banking dB**

3. 2-Phase Commit Protocol: Defining roles and assumptions

1. We define a **Transaction Coordinator TC**(either a separate dB node or an existing one): Master dB
2. We define **participants** of the commit process: 1 Master, 2 Slaves.
3. We assume there exists a **Write-Ahead-log** at each dB server.
4. Also we assume **dB server don't crash forever** and during a dB server crash the **Write-Ahead-log** is **NOT** lost/corrupted.



3. 2-Phase Commit Protocol

1. Let's say we want to perform a **withdraw transaction** onto a **banking dB(SQL) Customer Table** which has a **Master-Slave Replication Model**.

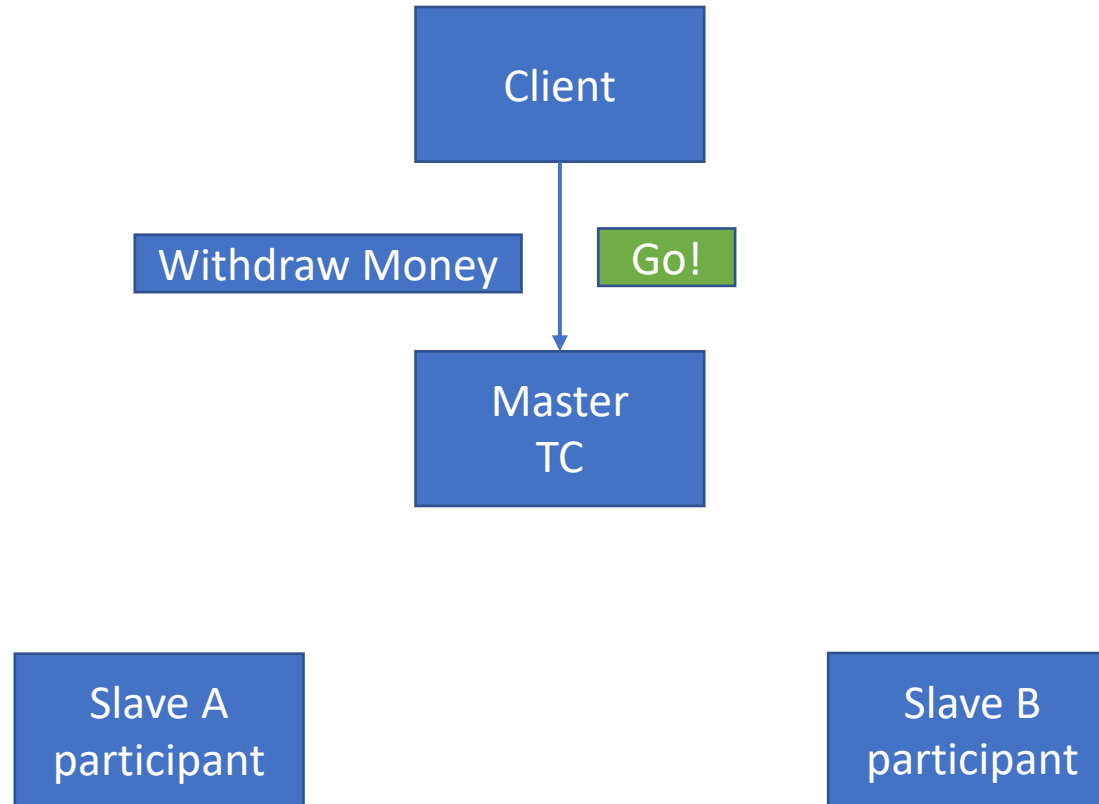
2. We have the following **withdraw transaction T**:

```
withdraw_money(Customer, amount) {  
    Begin_Transaction();  
    if (Customer.balance – amount >= 0) {  
        Customer.balance = Customer.balance – amount;  
        Commit_Transaction();  
    } else {  
        Abort_Transaction();  
    }  
}
```

3. 2-Phase Commit Protocol

1. Client(App Server) sends a “**withdraw_money**” request to Master dB(Transaction Coordinator)

C → TC: Go!

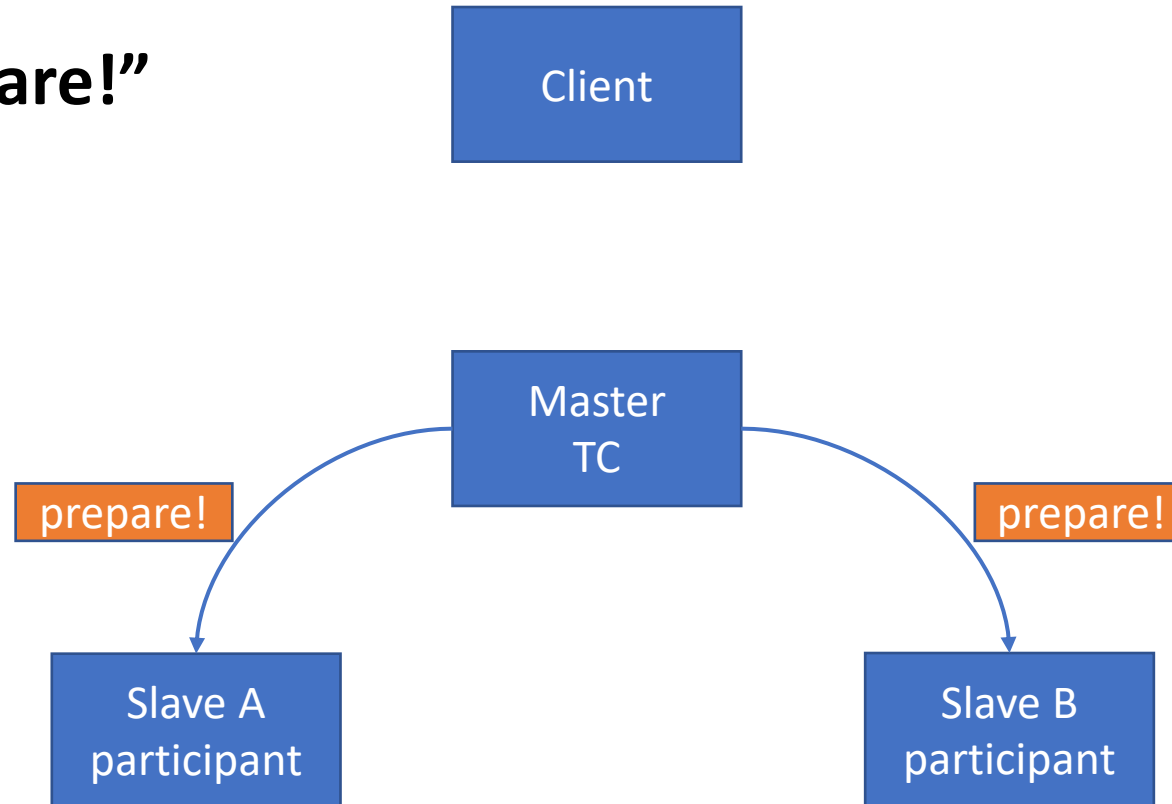


3. 2-Phase Commit: Voting Phase

2. Master(TC) sends a “**prepare!**” message to Slave A and Slave B.

TC→A: “**prepare!**”

TC→B: “**prepare!**”

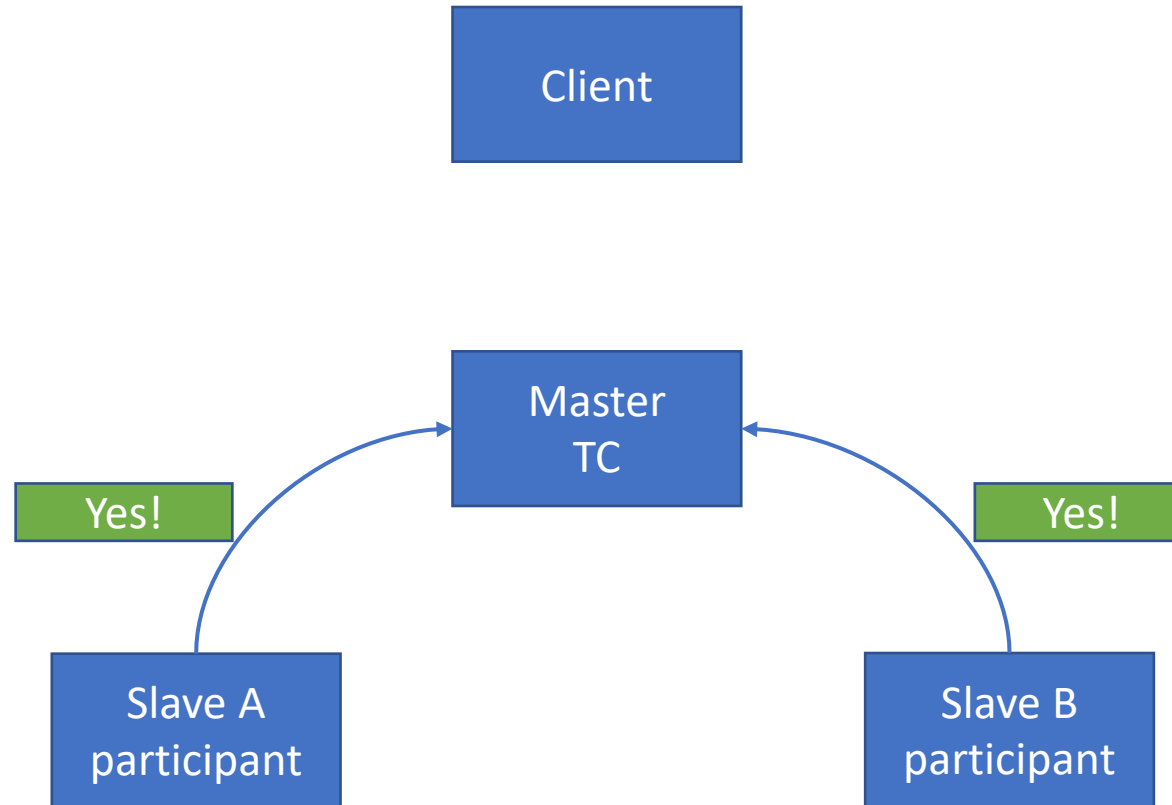


3. 2-Phase Commit: Voting Phase

3. Slave A and Slave B **executes** the transaction T upto before **commit_transaction()** and sends “**yes**” or a “**No**” to TC.

A → TC: yes!

B → TC: yes!

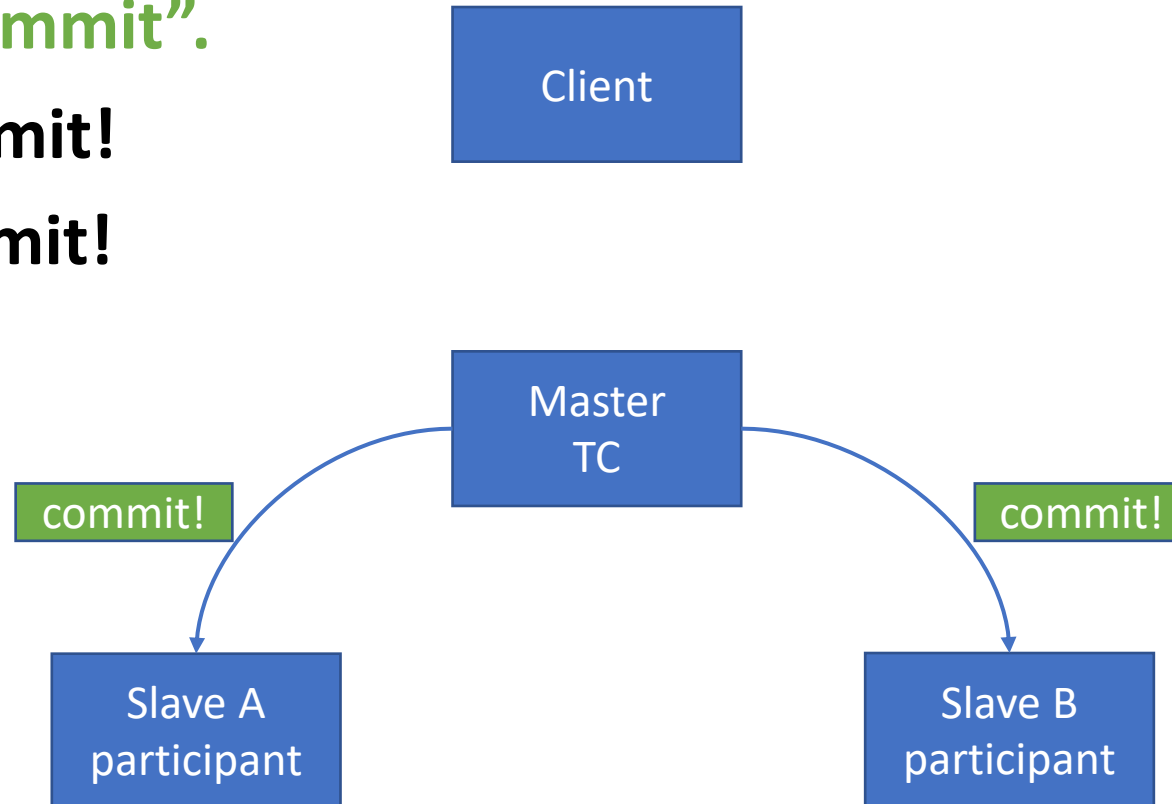


3. 2-Phase Commit: Commit Phase

4. TC sends a **commit** message if both A and B say “**yes**”. TC sends an **abort** message if either says “**No**”. Slave A and Slave B **commit** on receipt of “**commit**”.

TC → A: **commit!**

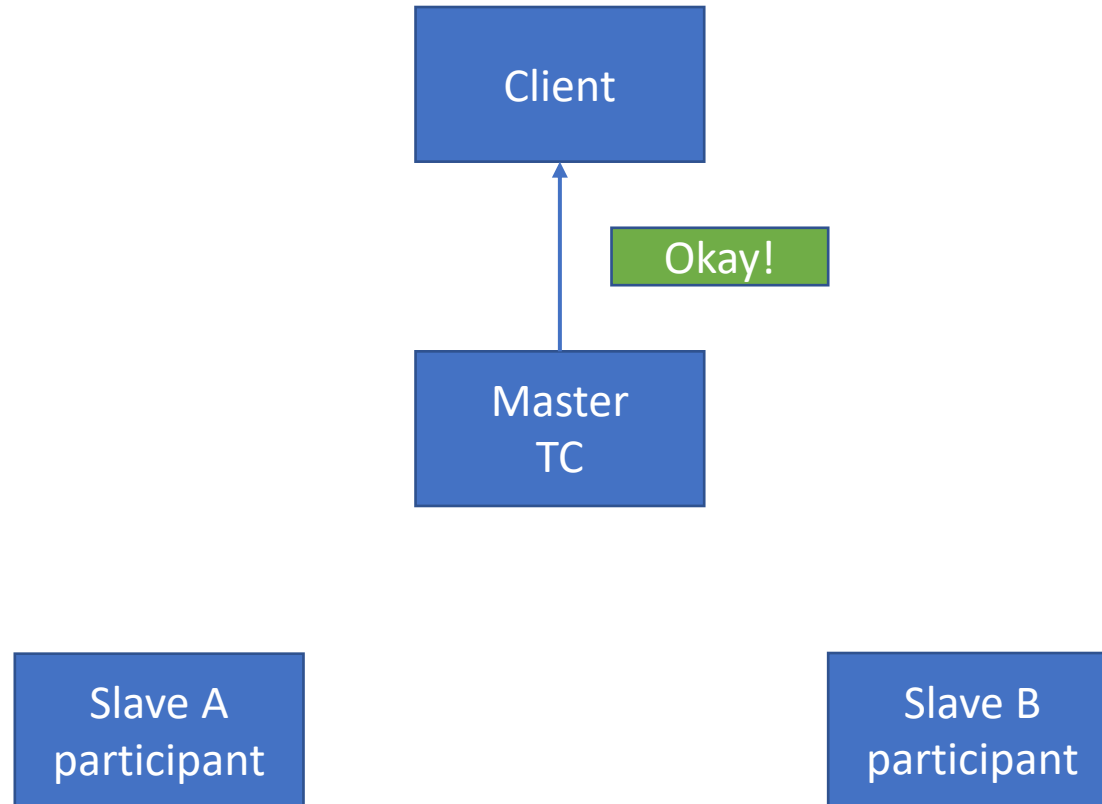
TC → B: **commit!**



3. 2-Phase Commit: Commit Phase

5. TC **commits** and sends “**Okay**” or “**Failed**” to the client.

TC → C: Okay!



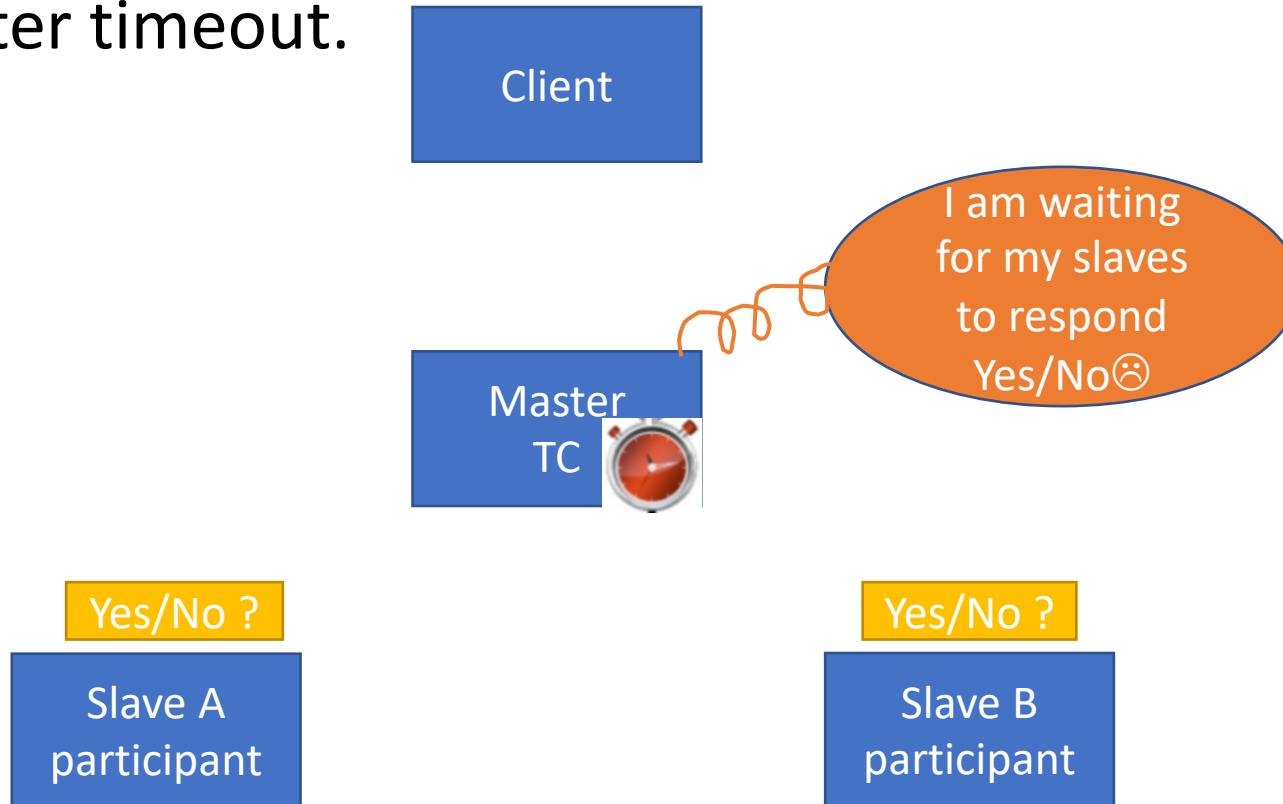
3. 2-Phase Commit: Atomic Commit Rule

1. If one dB **commits**, no dB **aborts**.
2. If one dB **aborts**, no dB **commits**.

So no dB can commit until **all the dBs agree together** to commit.

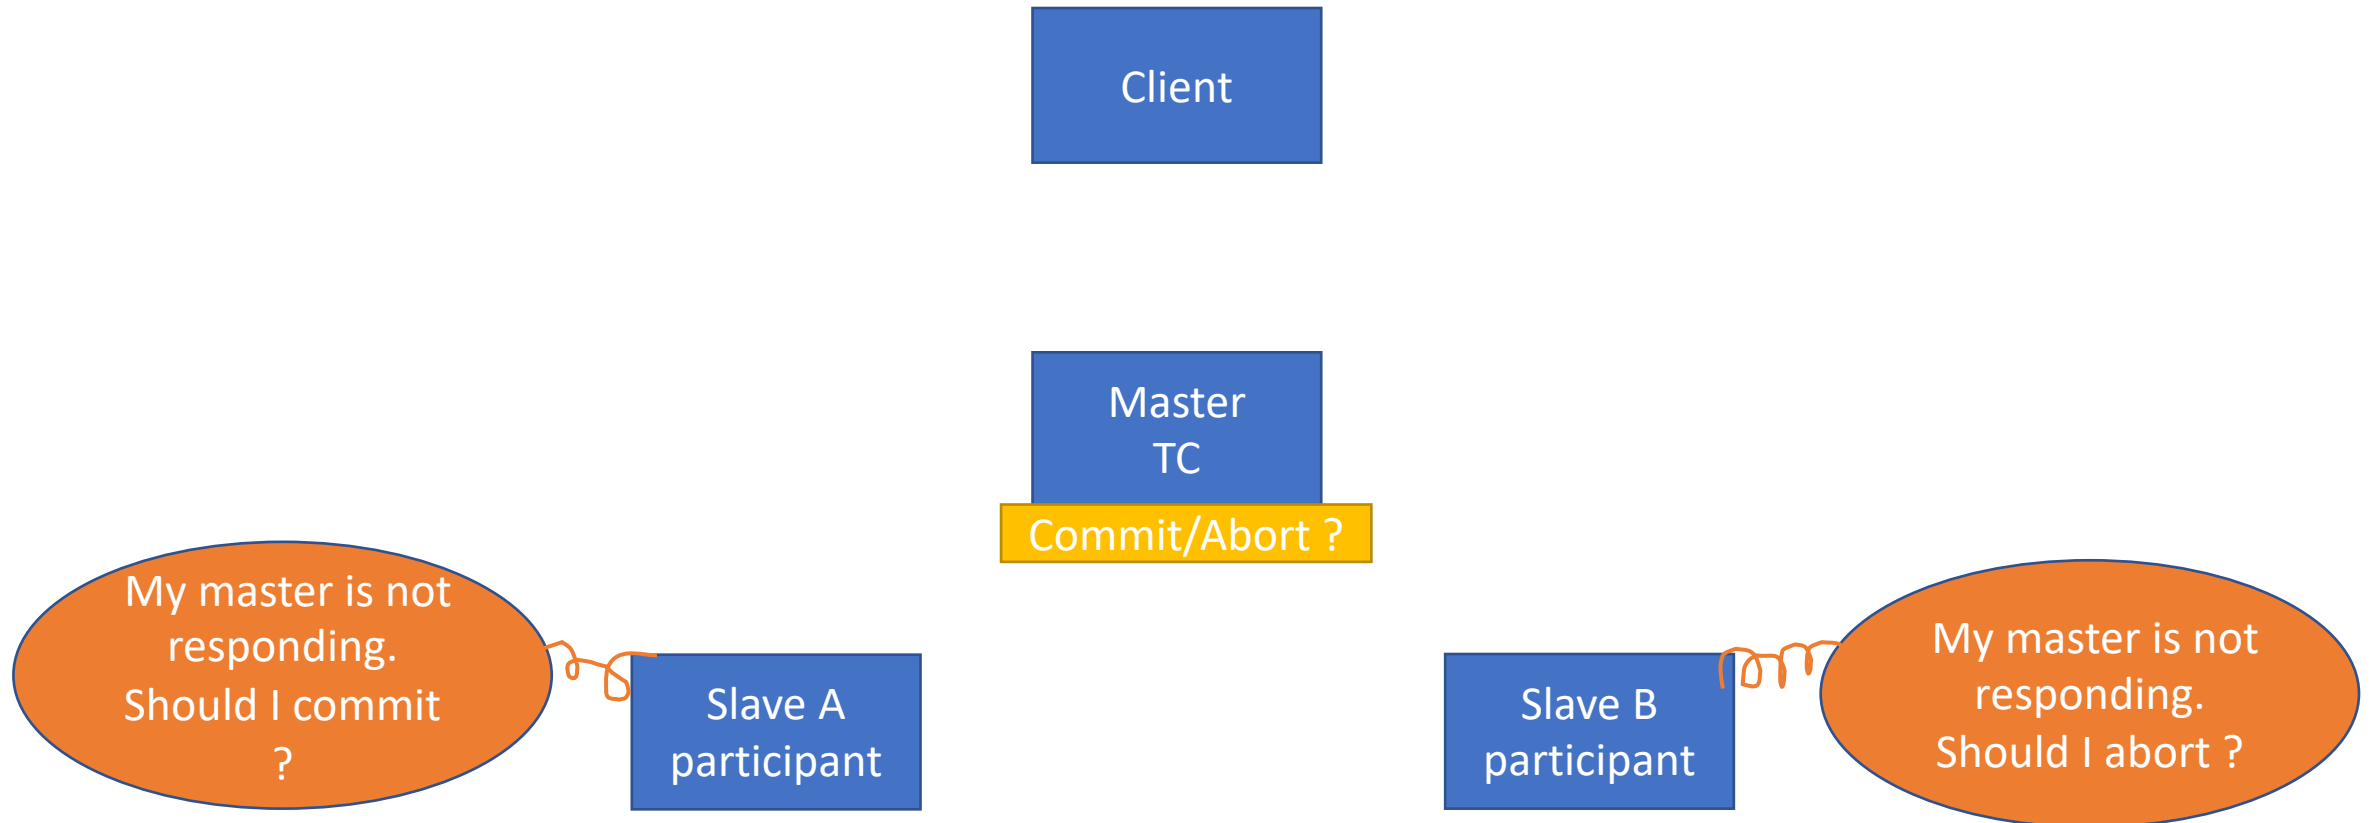
3. Timeouts in Atomic Commit

1. Let's say TC(Master) is waiting for a “**yes**” or “**no**” from A and B during **Voting Phase** ? TC hasn't sent any commit message yet so can safely **abort** after timeout.



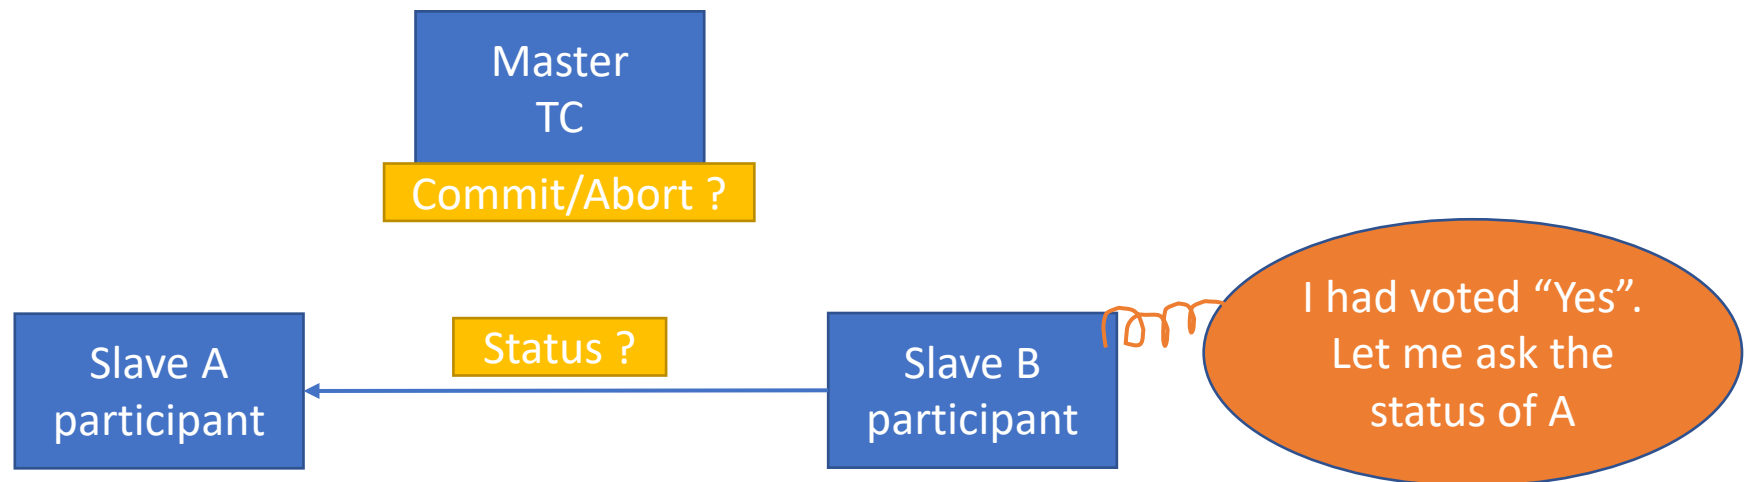
3. Timeouts in Atomic Commit

2. Let's say in **Commit Phase**, A and B is waiting for “**commit**” or “**Abort**” message from TC. **How long** should A and B wait ?



3. 2-Phase Commit: Termination Protocol

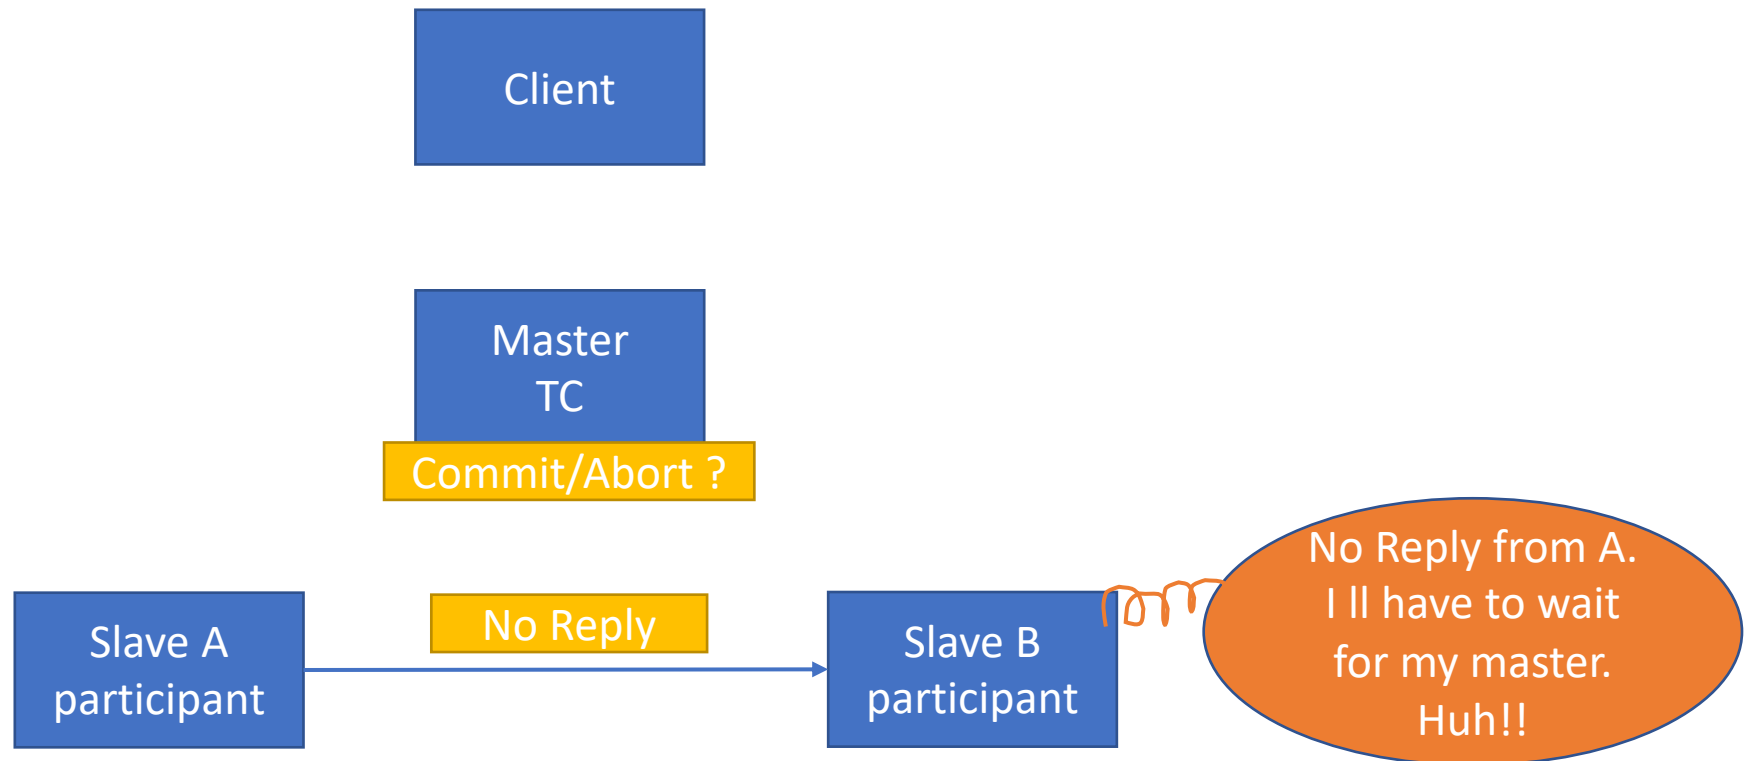
1. Let's say Slave A and Slave B are waiting for "**Commit**" or "**Abort**" message from TC(Master) in Commit Phase.
2. Let's assume Slave B had voted "Yes" in Voting Phase.
3. Slave B sends a message to Slave A asking for its status:
B→A: Status?



3. 2-Phase Commit: Termination Protocol

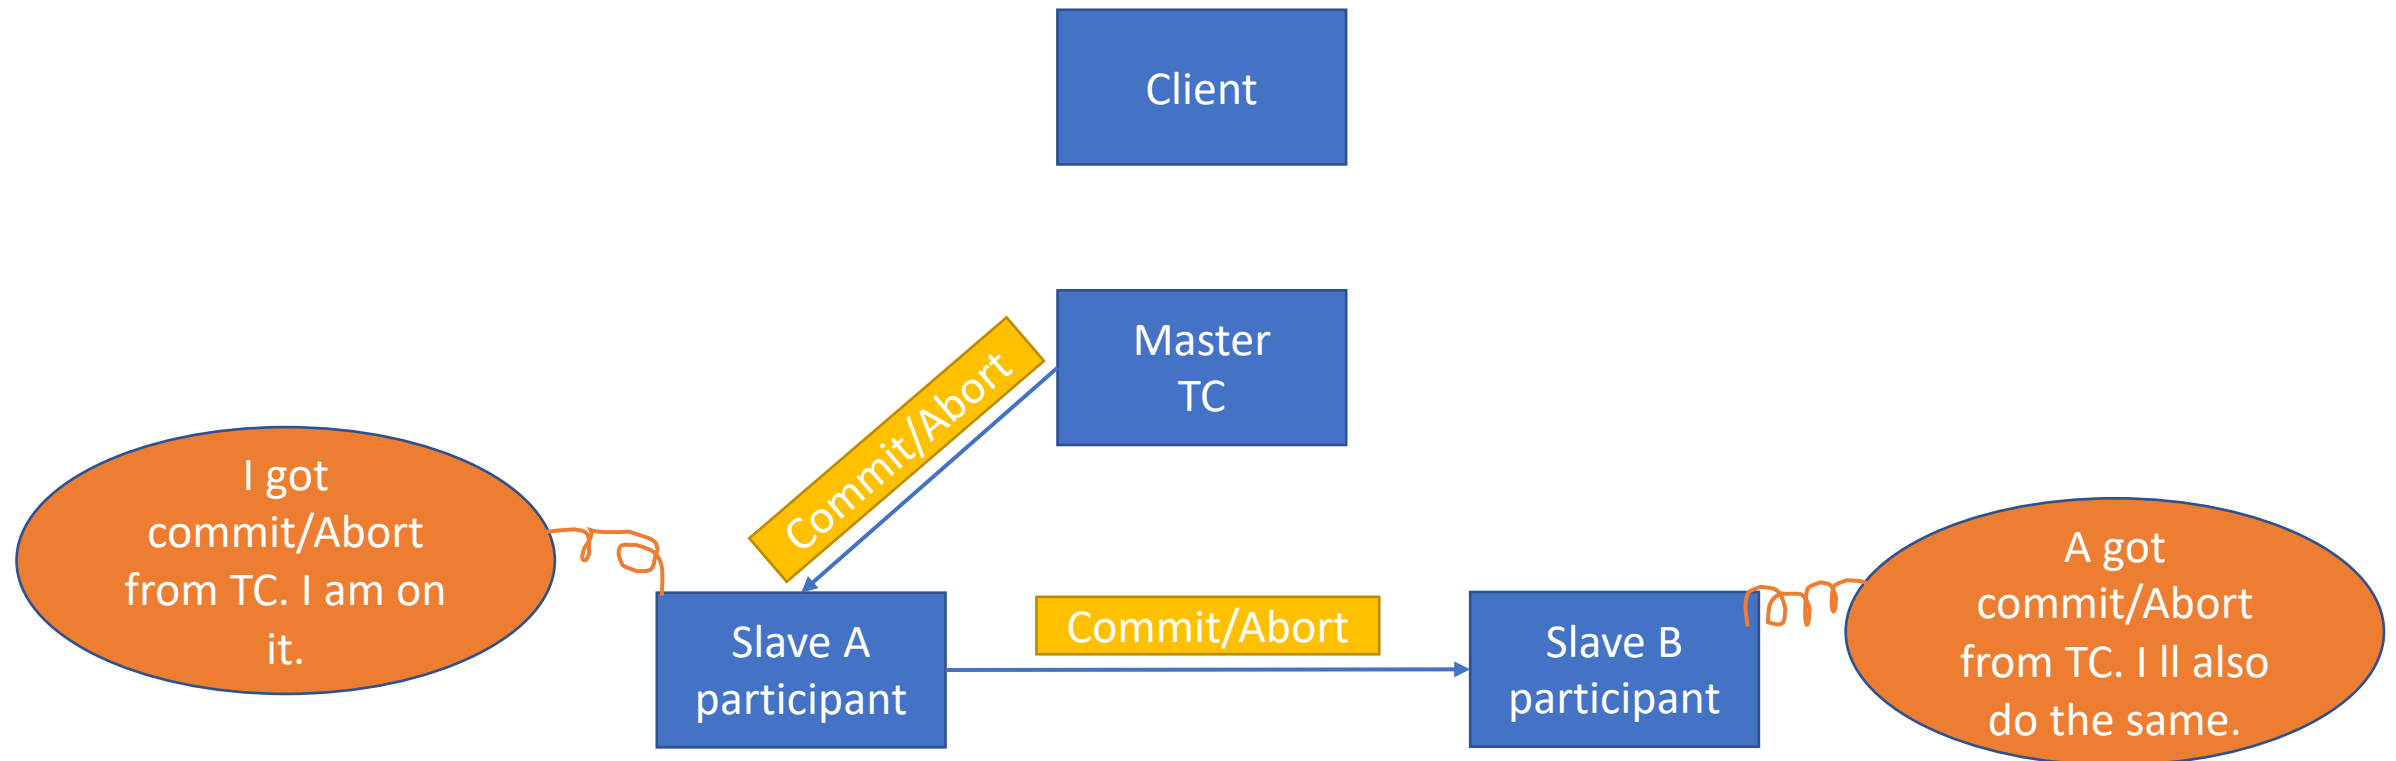
1. Four Cases arise, **Case 1:**

A → B: No Reply. B continues to wait for TC



3. 2-Phase Commit: Termination Protocol

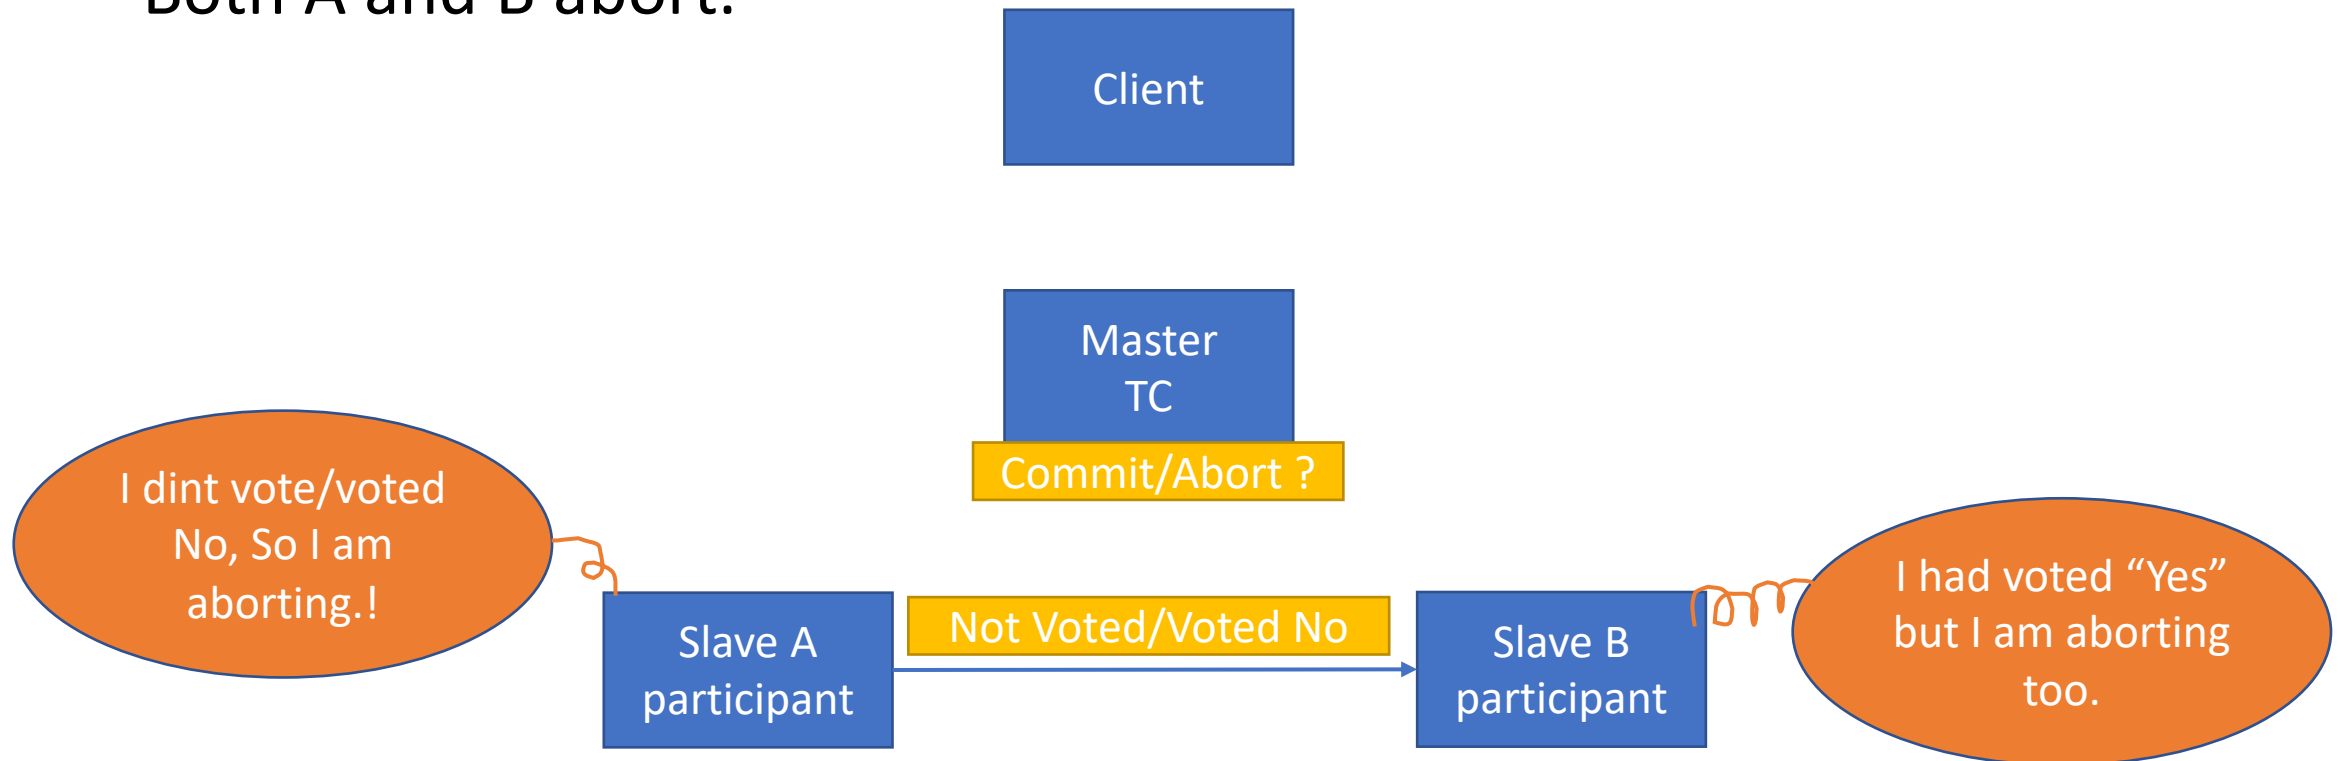
Case 2: $A \rightarrow B$: I got “commit”/”Abort” message from Master(TC).
B agrees with the Master’s(TC) decision.



3. 2-Phase Commit: Termination Protocol

Case 3: $A \rightarrow B$: I haven't voted yet or voted a "No". So TC can't have decided to "Commit".

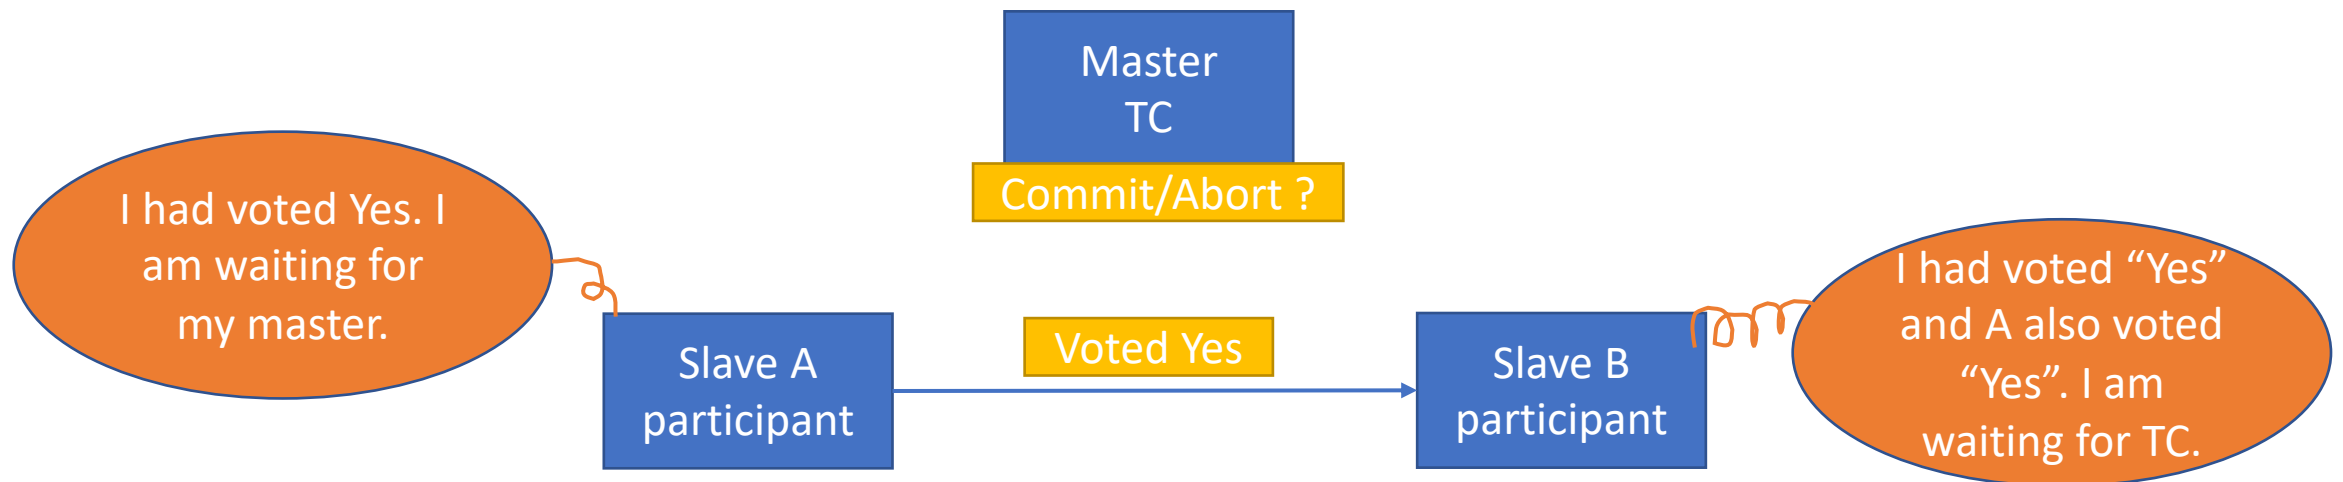
Both A and B abort.



3. 2-Phase Commit: Termination Protocol

Case 4: $A \rightarrow B$: I have voted “Yes”.

Both A and B must wait for TC. TC decides to “Commit” if both “Yes” received. TC “aborts” if timed-out.



3. 2-Phase Commit: Handling Crash/Reboot

1. What happens if **Master(TC) crashes** just after sending “**Commit**” in Commit Phase ?
2. What happens if Slave A or Slave B **crashes** just after sending “yes” during Voting Phase ?
3. All the dB nodes use their “**Write-Ahead-log**”(on disk) to record their state before crash/Reboot.
4. When every node **reboots** and is reachable, it follows **recovery Protocol** to perform commit/Abort.

3. 2-Phase Commit: Recovery Protocol

1. Master(TC) reboots and finds no “**Commit**” record on disk. It aborts.

TC: I did not send any commit message before so I am aborting.

2. A and B reboots and doesn't find “**Yes**” record on disk, It aborts.

A/B: I dint vote “Yes” so TC could'nt have committed.

3. A and B reboots and and finds “**Yes**” record on disk, it executes termination protocol.

3. Network Failure between dB Nodes

How does a dB system behave in case of network failure between their nodes ?





4. CAP Theorem

4. CAP Theorem

It says that it is **IMPOSSIBLE** for a distributed dB system to achieve all the three Consistency, Availability and Partition Tolerance. We can pick **only two** of them.

1. **Consistency:** Every dB read request receives the value of **most recent write** or an **error**.
2. **Availability:** Every dB request receives a **non-error response**. There is no guarantee that the response contains the most recent writes(uptime = 99.99% of total time of service).
3. **Partition Tolerance:** The dB system continues to operate **even if any number of messages are dropped/delayed** by network between dB nodes.

4. Consistency Patterns

1. **Weak Consistency:** After a write, reads may or may or may not see it. E.g – App Engine: memcache, VoIP, live online vid, Realtime multiplayer game
2. **Eventual Consistency:** After a write, reads will eventually see it, mail, Search Engine, Indexing, DNS, Amazon S3
3. **Strong Consistency:** After a write, reads WILL see it, App Engine: Data Store, File System, RDBMS, Azure Tables.

4. Availability Numbers

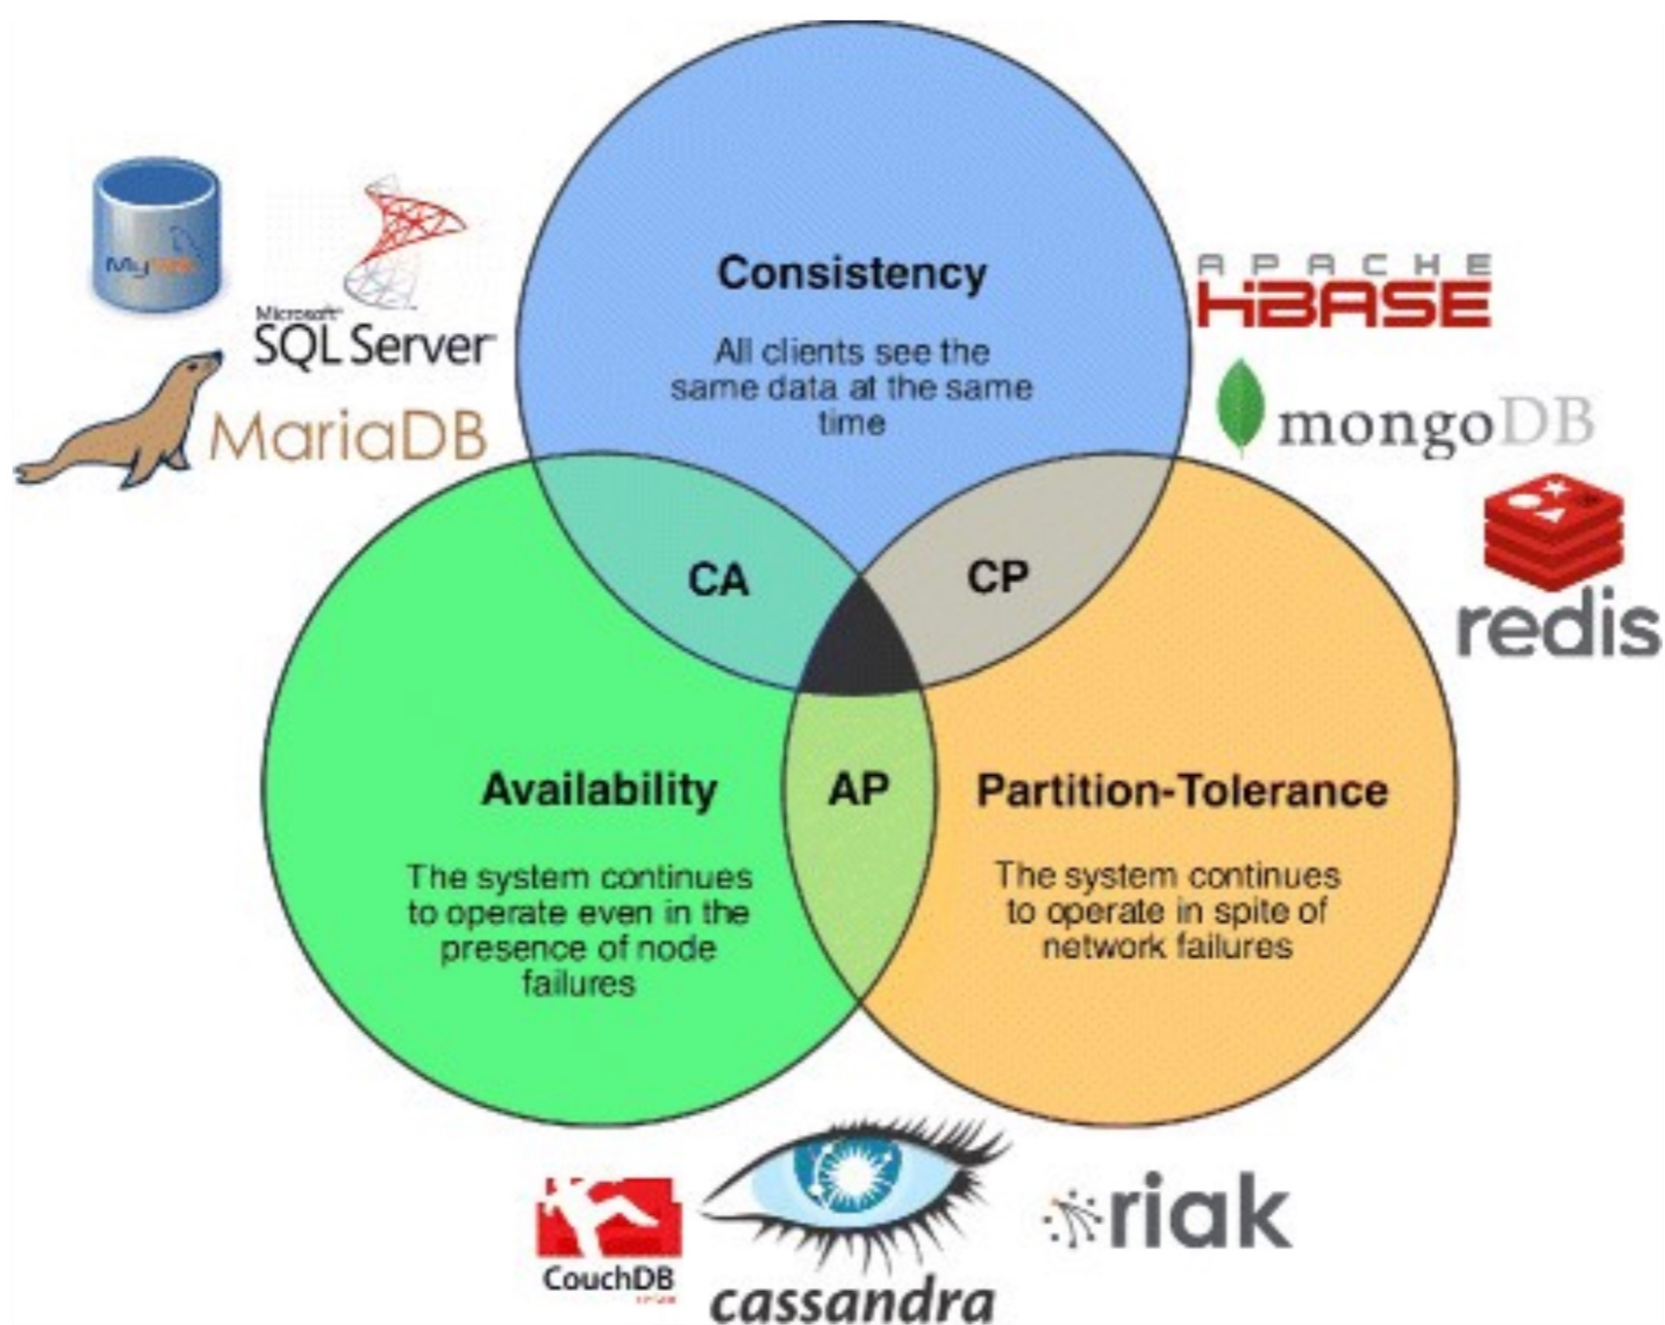
Availability is the **Uptime** of a system as a percentage of **total time of service**. It is measured in number of 9s.

99.9% availability - three 9s

Duration	Acceptable downtime
Downtime per year	8h 45min 57s
Downtime per month	43m 49.7s
Downtime per week	10m 4.8s
Downtime per day	1m 26.4s

99.99% availability - four 9s

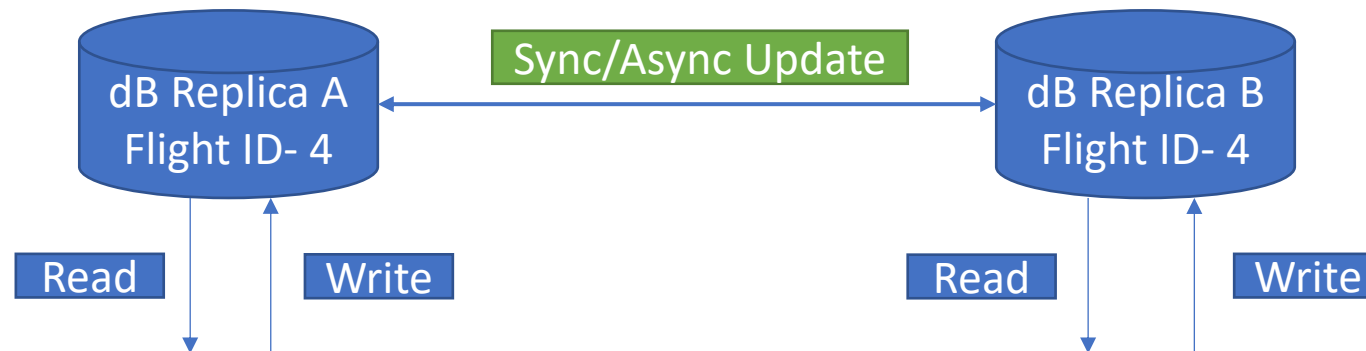
Duration	Acceptable downtime
Downtime per year	52min 35.7s
Downtime per month	4m 23s
Downtime per week	1m 5s
Downtime per day	8.6s



CAP Theorem

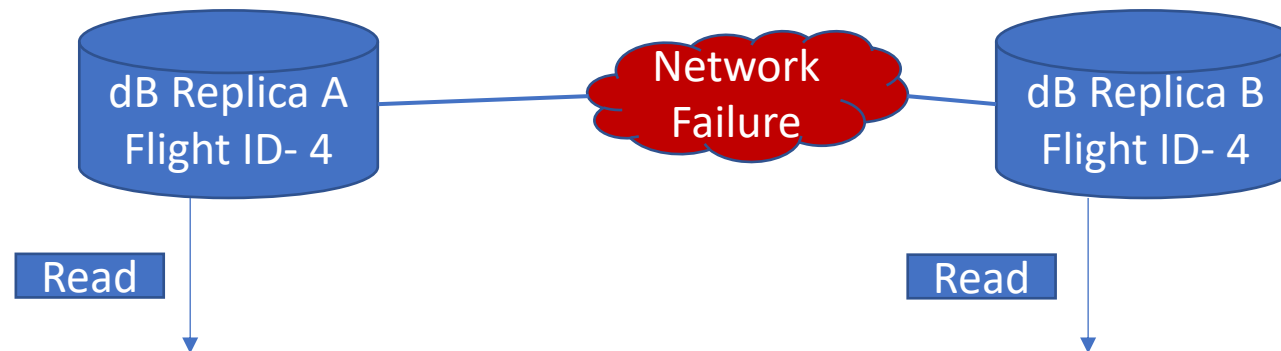
4. CAP Theorem: CA

1. A **write transaction** happening at dB Replica A will be updated to dB Replica B and vice versa. Both replica **maintain consistent copies** of data at any point of time. So every read receives most recent write.
2. Also dB A and B is **up(available)** for read/Write 99.99% of total time. So this provides CA.



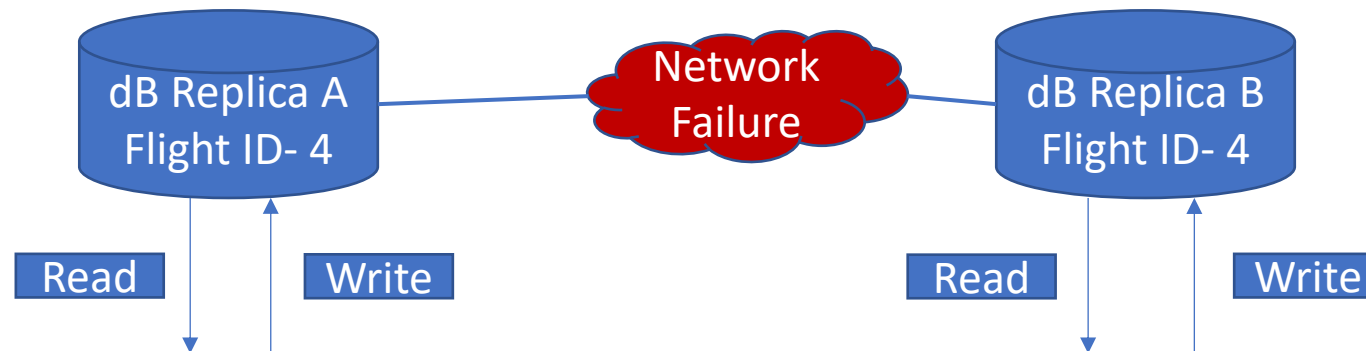
4. CAP Theorem: CP[SQL dBs]

1. Lets say the **network** connecting Replica A and B goes **down** and dBs get partitioned. Now both dBs stop serving Write requests as network failure is detected. The system is **not available** but maintains **consistent copies** of data in Replica A and B.



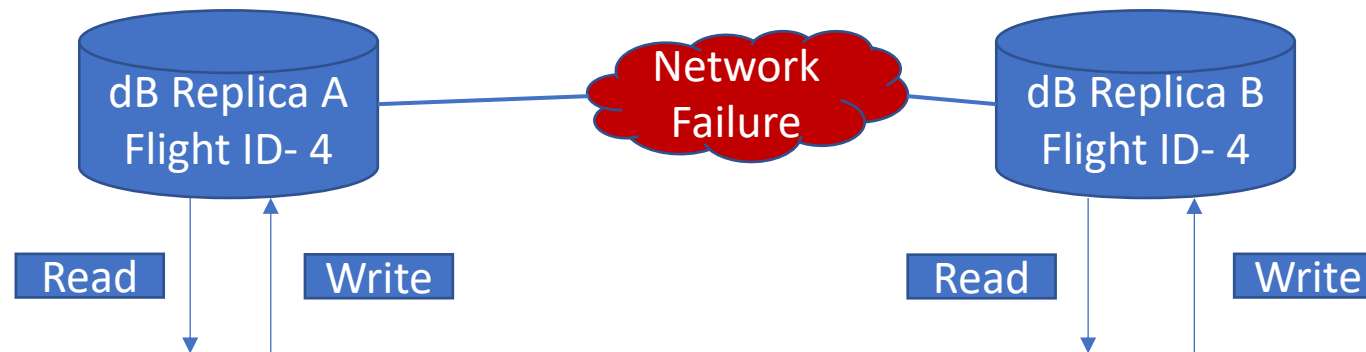
4. CAP Theorem: AP[NoSQL dBs]

1. Lets say the **network** connecting Replica A and B goes **down** and dBs get partitioned. Both dB Replicas are running fine.
2. Now a write transaction happening at A/B **can no longer** reach dB B/A respectively. So both dB A and B end up maintaining **inconsistent copies** of data. But it continues to **serve(Available)** in a disconnected env.



4. AP and Eventual Consistency

1. The idea is the dB system **will become consistent over time** after **network failure is fixed**.
2. We **allow dB read/Write** requests on both dB Replicas even after network failure.

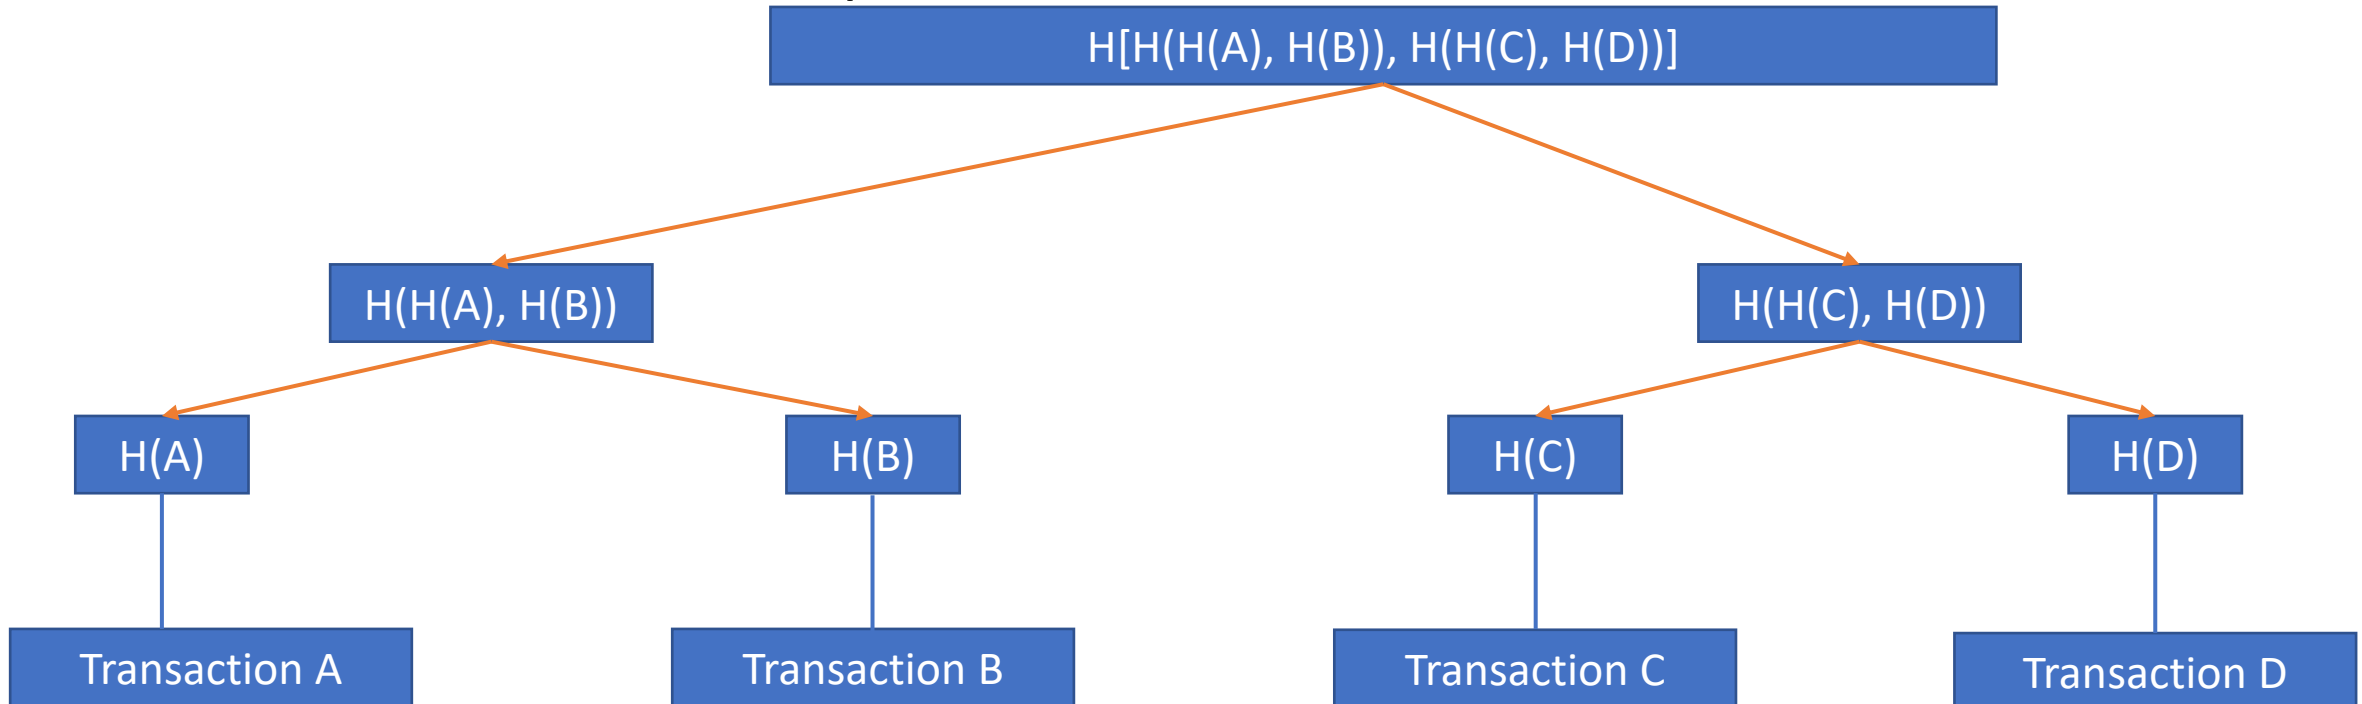


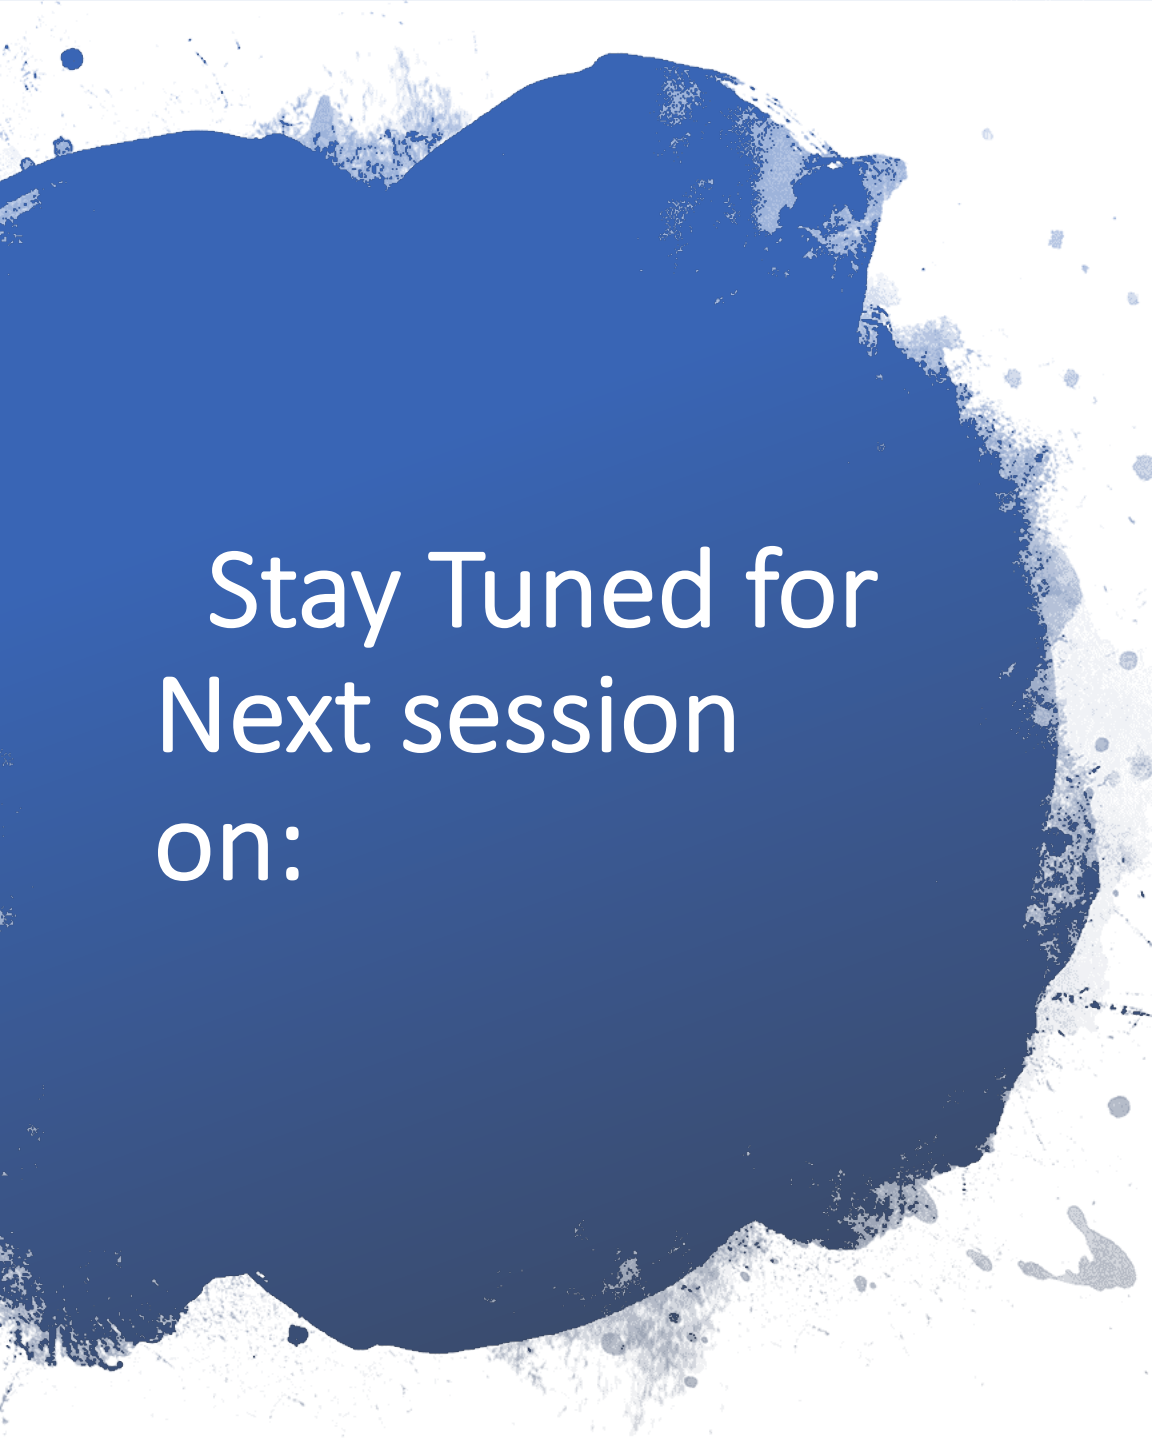
4. AP and Eventual Consistency

1. Both dB Replica A and B maintains a **Merkle Tree** which has the **latest Write Transactions**.
2. Lets say 4 write transactions A, B, C, D occurred on dB A during network Failure.
3. Each dB Replica maintains a **Merkle Tree of write transactions**.

4. Merkle Tree

1. We hash the transaction from leaf up to the root of tree **at each dB**.
2. Now we perform **Tree Traversal($O(N)$)** and do **data sync between dBs** as network is up.





Stay Tuned for
Next session
on:

Caching and CDN



Questions

